

# L6

# APIs and DB Security

Eugene Wu

# Topics

Interfacing with applications

- Types of DBMSes

- Database APIs (DBAPIS)

- Cursors

Security

- SQL injection, the world's most popular DB “hack”

- Access controls and GRANT

- Encryption

# How to use SQL in an application?

SQL is not a general purpose programming language

Designed for data access/transform, and optimization

Applications are complex, require “business logic”

# Many Options

Extend SQL to be turing complete

- Makes optimization very very hard
- Technically, `recursive WITH` clause makes SQL turing complete...

Embedded SQL

- extend language by “embedding” SQL into it

DBAPI

- Low-level library with core database calls

Object-relational mapping (ORM)

- Define DB-backed classes, magically map between objects & DB tuples

# Embedded SQL


Extend host language (python) with SQL syntax

e.g., `EXEC SQL sql-query`

goes through a preprocessor

Compiled into program that interacts with DBMS directly

# Embedded SQL



```
...  
if (user == 'admin'){  
    EXEC SQL select * ...  
} else {  
    ...
```

# What does a library need to do?


Single interface to possibly multiple **DBMS engines**

Connect to a database

Manage transactions (later)

Map objects between host language and DBMS

Manage query results



# DBMS Engines



Web app


Data Science

ML

Declarative Interface (SQL)


DBMS

Logical Schema



Logical  
Query Plan

Query  
Optimizer



Physical  
Query Plan

Plan  
Executor


Recovery &  
Concurrency Control

Files and Access Methods

Buffer Pool to manage memory

Memory

Disk Storage



Web app

Data Science


ML

Declarative Interface (SQL)

DBMS

DBMS Engine

Disk Storage



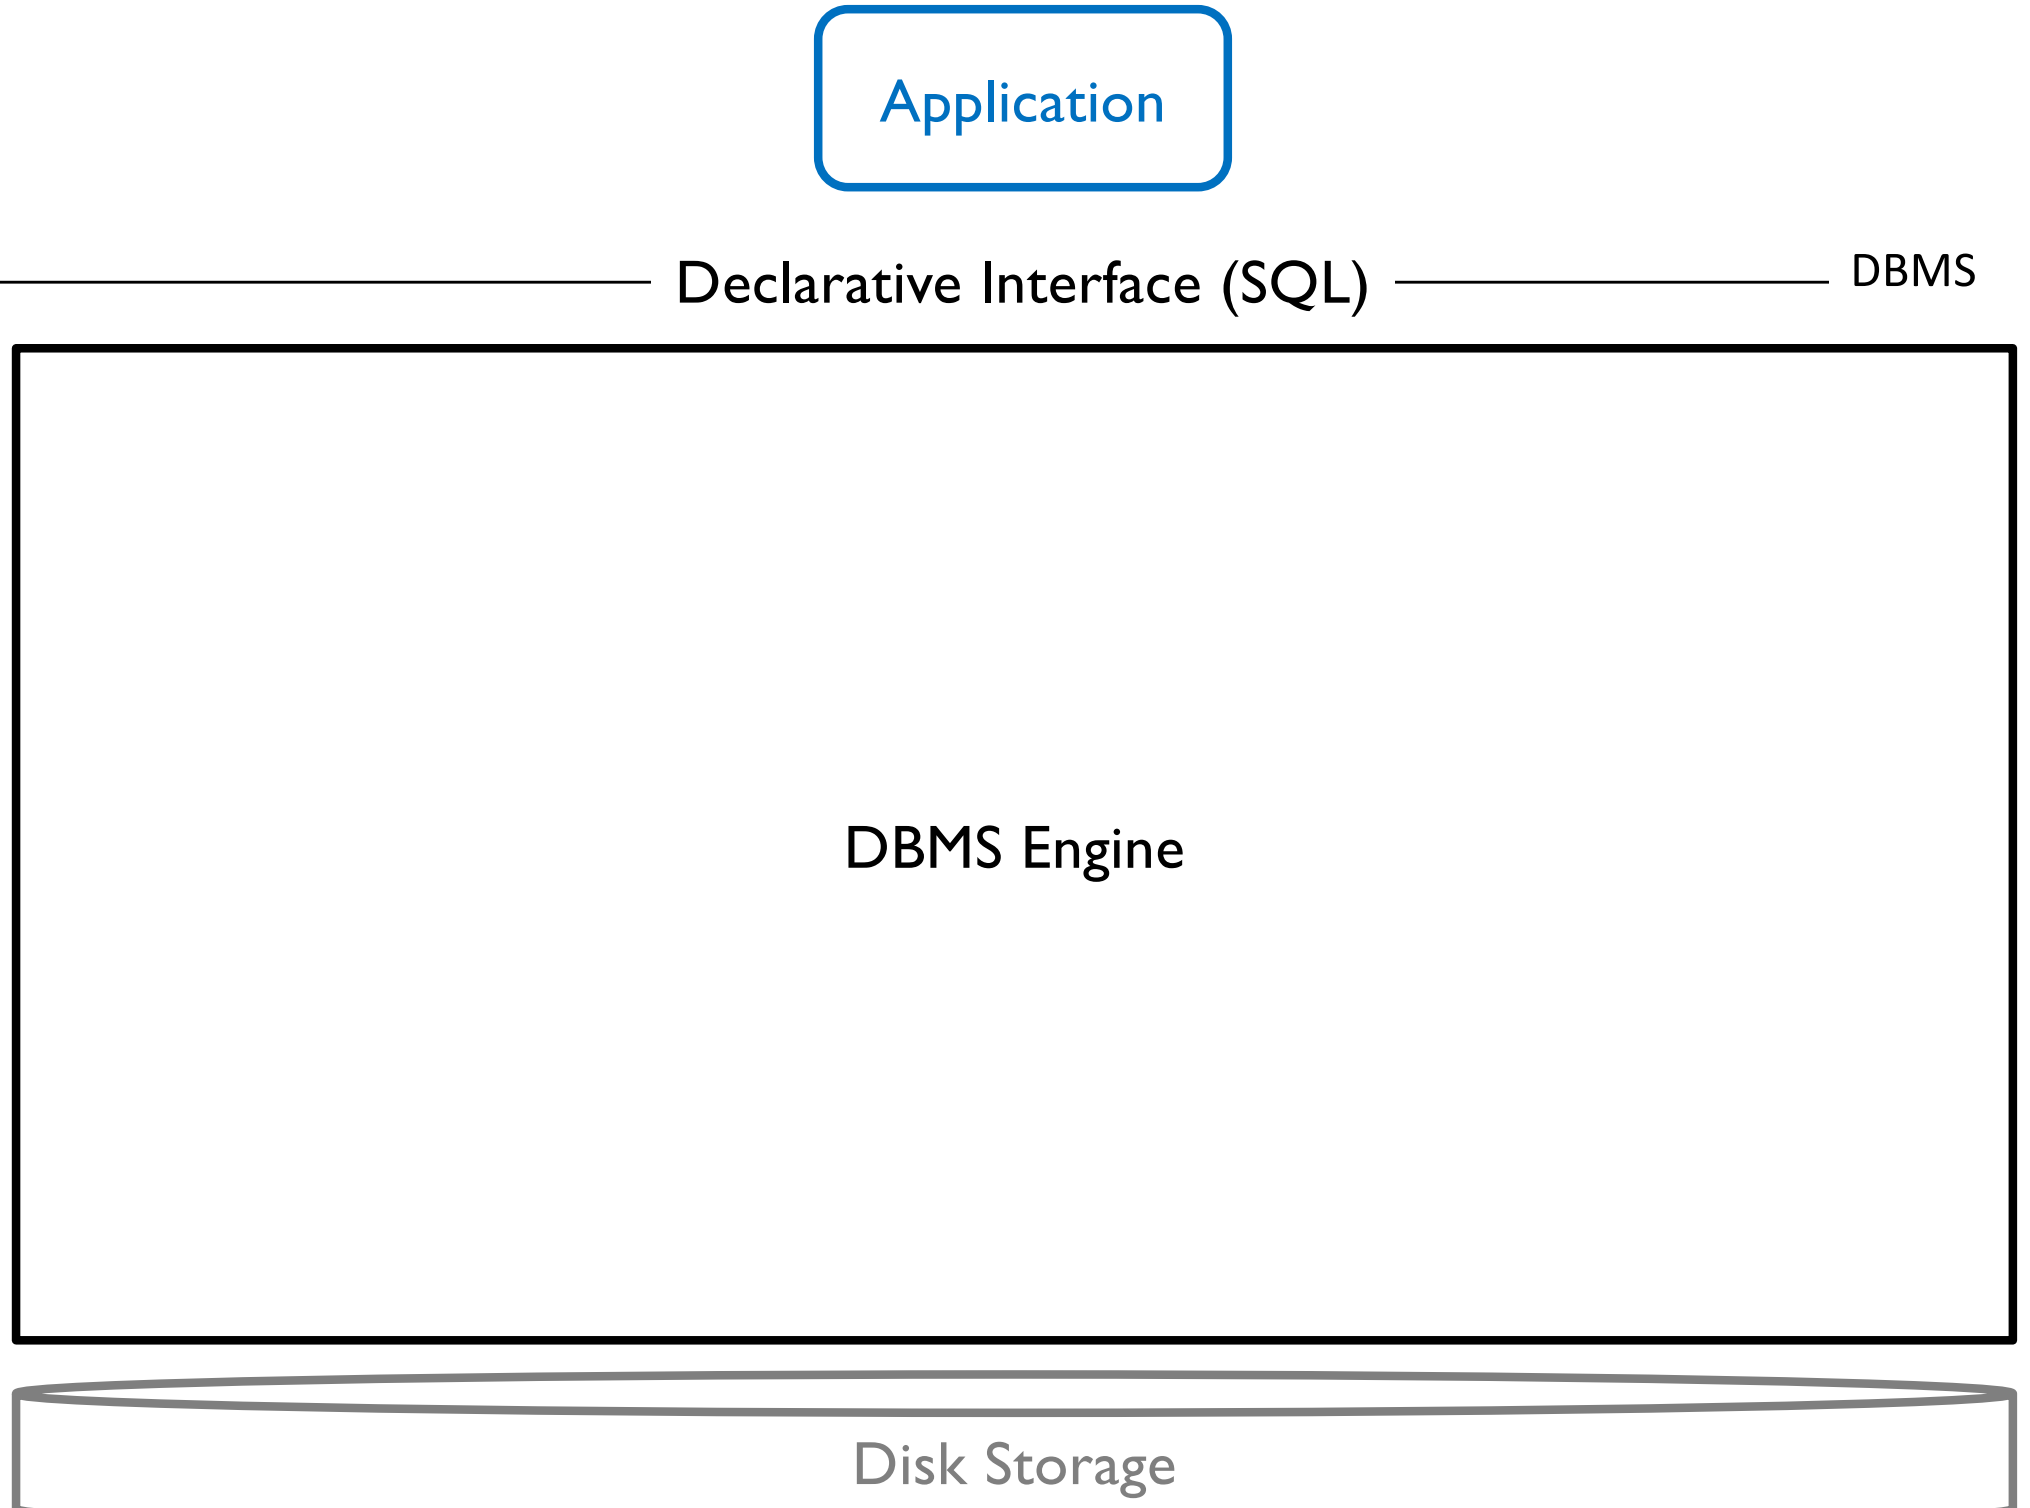
Application

Declarative Interface (SQL)

DBMS

DBMS Engine

Disk Storage





The diagram illustrates the layers of a database system. At the top is the 'Application' layer, represented by a blue-outlined rounded rectangle. Below it is the 'Declarative Interface (SQL)' layer, indicated by a horizontal line. This line connects to the 'DBMS' label on the right and the 'DBMS Engine' box in the center. The 'DBMS Engine' is a black-outlined rectangle. At the bottom is the 'Disk Storage' layer, represented by a wide, shallow gray-outlined oval.

Application

Declarative Interface (SQL)

DBMS

DBMS Engine

Disk Storage

# DBMS Engines




```
graph TD; Application[Application] --- DBMS_Engine[DBMS Engine];
```

Application

DBMS Engine

# DBMS Engines: 3 Types




# Client-Server DBMS on Different Machines

Main DBMS logic runs on server process

Apps use DBMS client library to connect with server


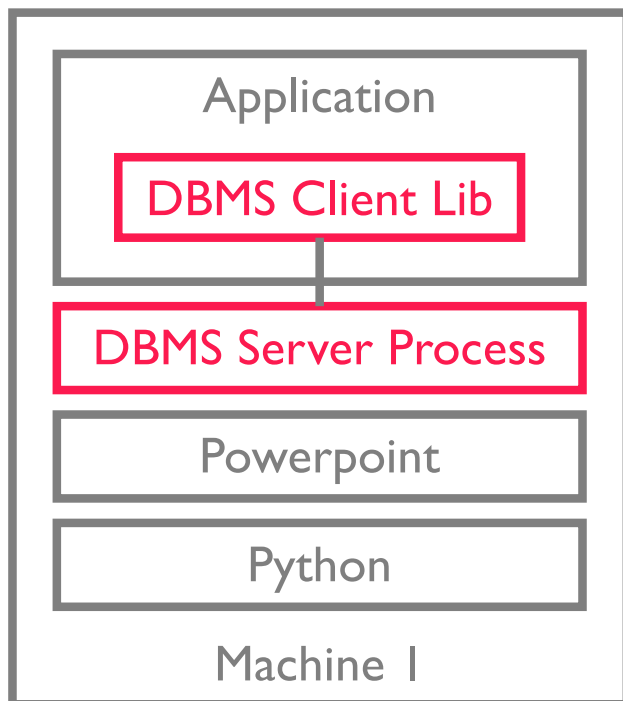
Usually communicate via a network protocol such as TCP



# Client-Server DBMS on Same Machine

Server process can run on the same machine as application

Usually communicate via TCP, Interprocess Communication (IPC)



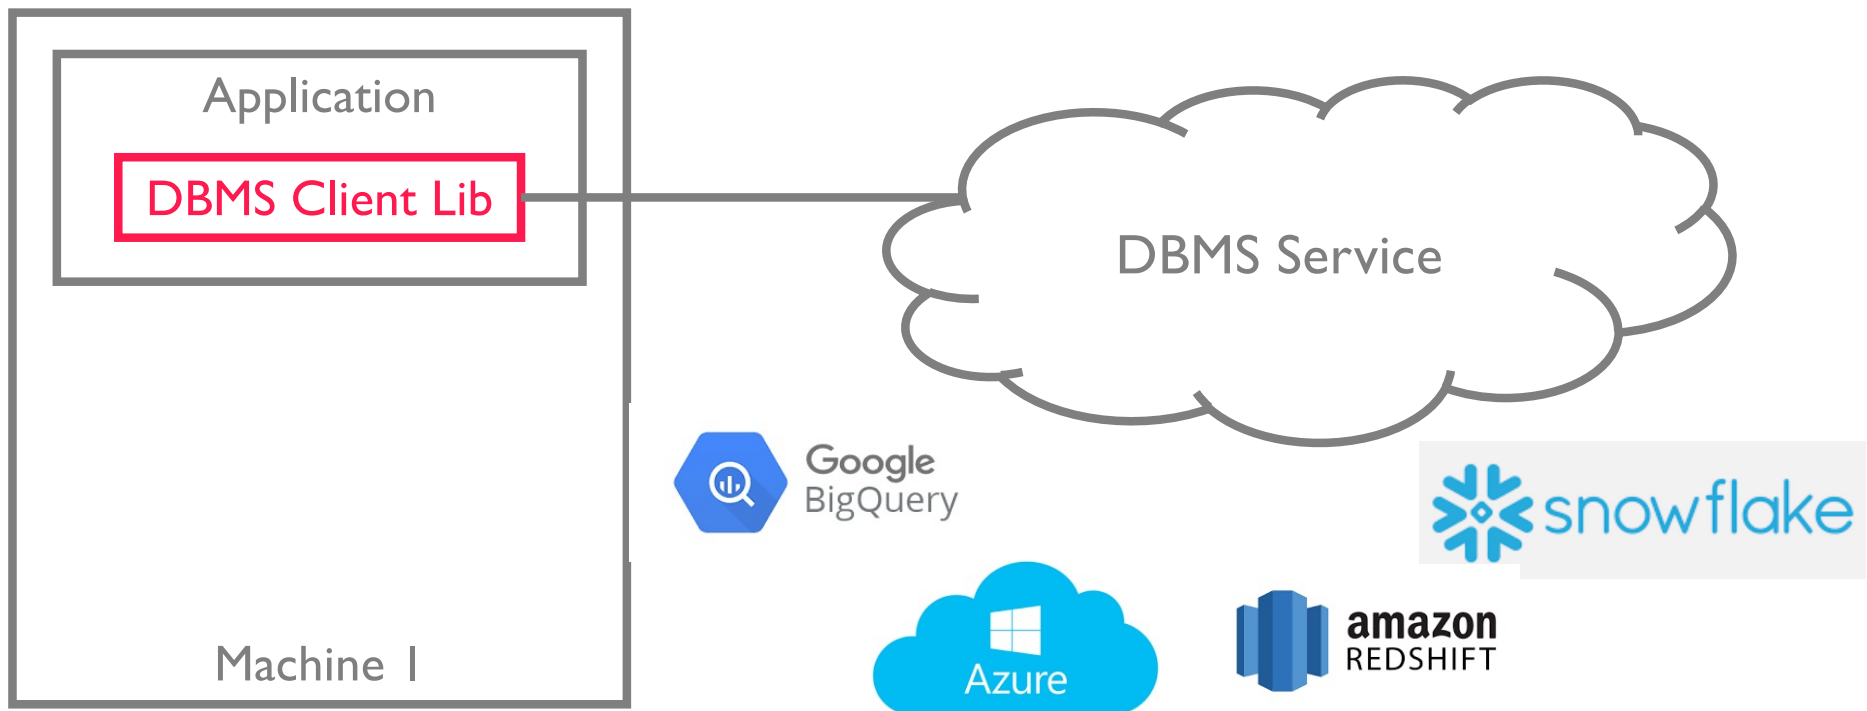


# Cloud DBMS

DBMS service managed by someone else

Meant to be auto-scaling (add/remove machines based on load)

Communicate via network protocol e.g., TCP/HTTPS



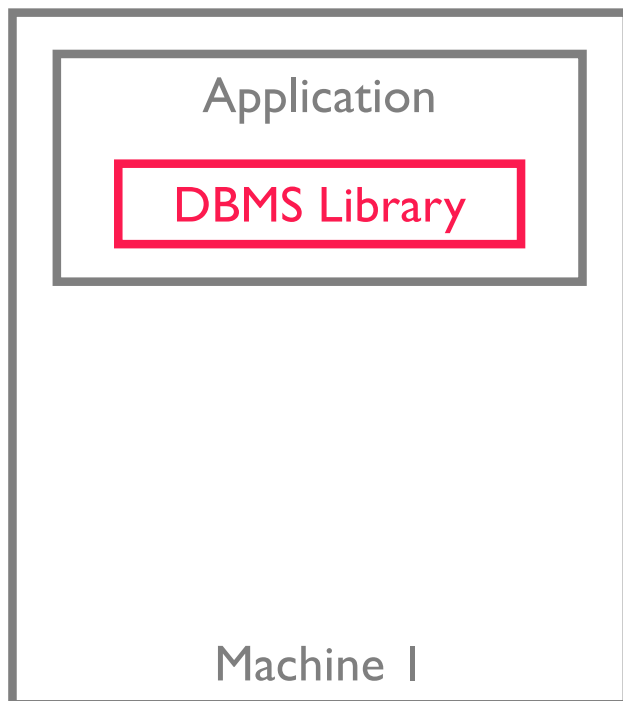
# Embedded DBMS

DBMS is a library linked by the application

```
import duckdb, sqlite
```

Runs in same process and memory space

No communication, usually in-memory

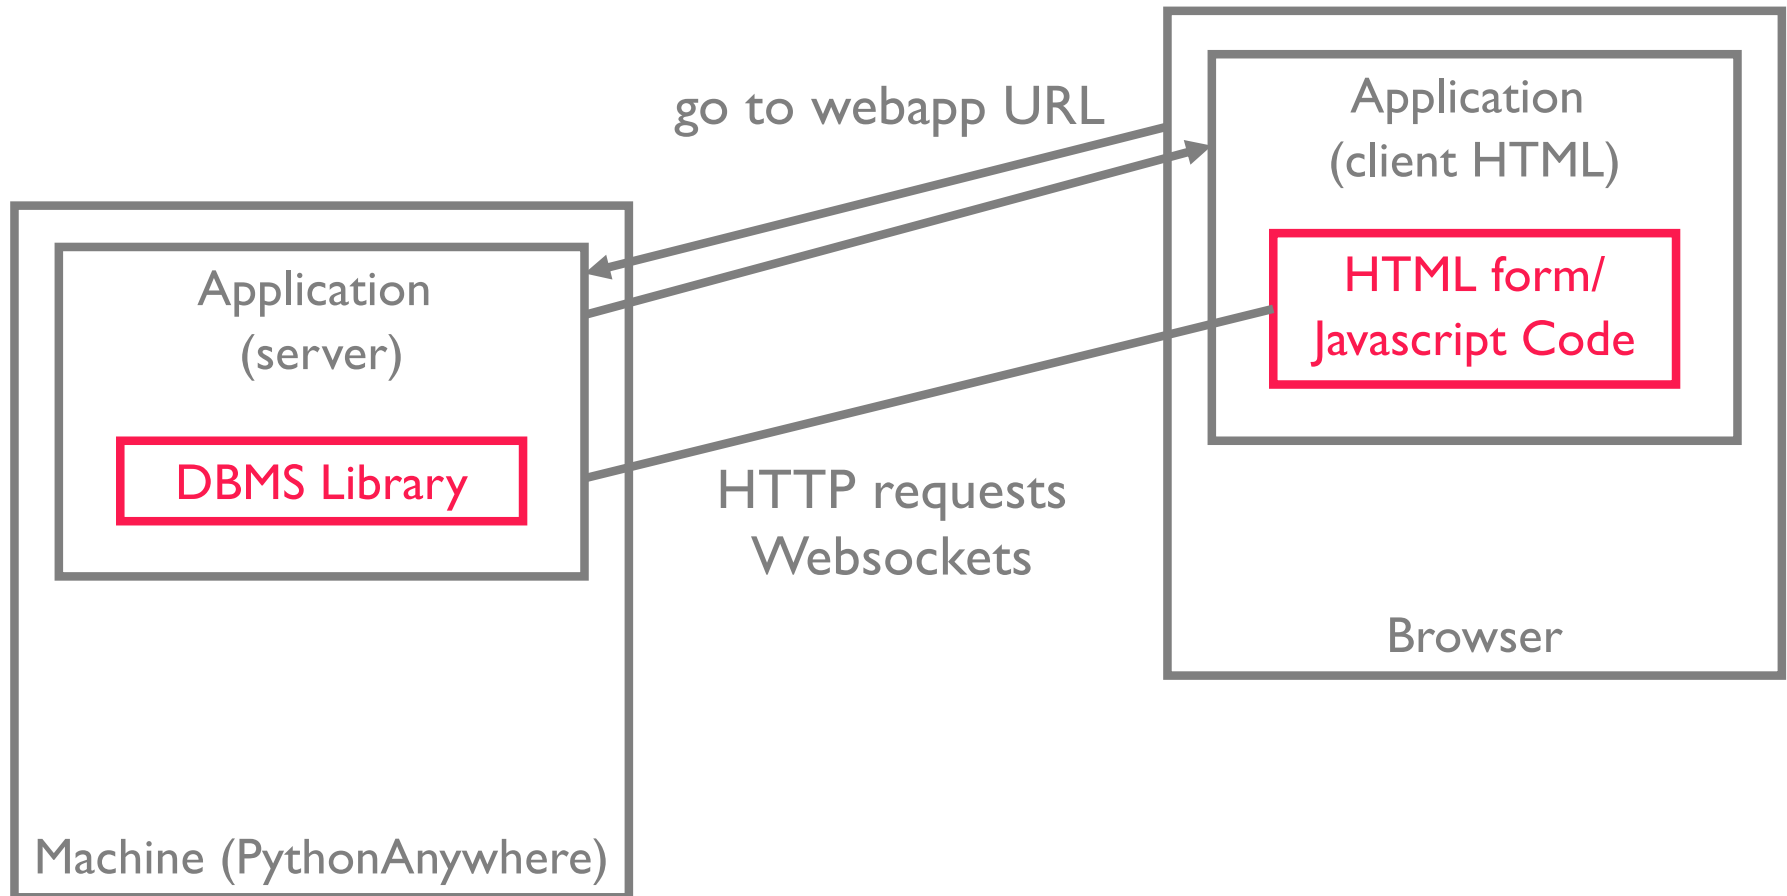


SQLite, DuckDB compiled to javascript!  
Can run directly in webpage.

# Applications in Practice

Web apps usually also have a client (browser) and server.

App's server *uses* DBMS client lib to connect to DBMS Server.



# Applications in Practice

## **Server code for `http://.../query`**

```
@app.route("/query/")
def query():
    name = request.form['name']
    cur = conn.execute("...")
    row = cur.fetchone()

    data = dict(val=row['attr'])
    return render_template(
        "results.html",
        data=data)
```

Receive inputs &  
issue query

Package query results  
into http response

# Applications in Practice

## Server code for `http://.../query`

```
@app.route("/query/")
def query():
    name = request.form['name']
    cur = conn.execute("...")
    row = cur.fetchone()

    data = dict(val=row['attr'])
    return render_template(
        "results.html",
        data=data)
```

## results.html

```
<div>
    result is {{data}}
</div>
```

Replace placeholder  
in template with data

# DB API Overview

## Library Concerns

1. Establish connection
2. Submit queries/transactions
3. Retrieve results

## Impedance Mismatches

1. Types
2. Classes/objects
3. Result sets
4. Functions
5. Constraints

# DB API: Engines

URI to refer to a given DBMS engine (like a URL)

`driver://username:password@host:port/database`

```
from sqlalchemy import create_engine
uri1 = "postgresql://localhost:5432/testdb"
# embedded dbmses have no host:port
uri2 = "sqlite:///testdb.db"

engine1 = create_engine(uri1)
engine2 = create_engine(uri2)
```

# DB API: Connections

Connect to the DBMS engine

- DBMS Server allocates resources for connection
- Relatively expensive, libs often cache+reuse connections
- Defines scope of a transaction (later in semester)

```
conn1 = engine1.connect()  
conn2 = engine2.connect()
```

Close connections when done to avoid leaking resources.

```
conn1.close()
```



# DB API: Query Execution

```
conn1.execute("UPDATE TABLE test SET a = 1")  
conn1.execute("UPDATE TABLE test SET s = 'wu'")
```

```
# sqlite
```

```
conn1.execute("SELECT * FROM test WHERE a = ?", 1)
```

```
# postgres
```

```
conn1.execute("SELECT * FROM test WHERE a = %s", 1)
```

# DB API: Query Execution

```
foo = conn1.execute("select * from big_table")
```

## Challenges

- Impedance mismatches

- What is the return type of `execute()`?

- How to pass data between DBMS and host language?

- Can we only pass data between DBMS and host language?

# (Type) Impedance Mismatch

SQL standard maps between SQL and several languages

Most libraries support primitive types

SQL types	C types	Python types
CHAR(20)	char[20]	str
INTEGER	int	int
SMALLINT	short	int
REAL	float	float

What about complex objects { x:'I', y:'hello' }

# (Class) Impedance Mismatch

Programming languages usually have classes

Want objects to persist in DBMS

```
user.name = "Dr Seuss"  
user.job  = "writer"  
  
class User { ... }  
class Employee extends User { ... }  
class Salaries {  
    Employee worker;  
    ...  
}
```

Object Relational Mappings (ORMs)  
try to provide this abstraction

# ORM: classes that magically sync with DBMS

Base is a special class defined by ORM

mimics CREATE TABLE in Python

We will NOT use ORMs for project I

```
class User(Base):  
    __tablename__ = "user_account"  
    id            = Column(Integer, primary_key=True)  
    name         = Column(String)  
    addrs        = relationship(  
        "Address", back_populates="user", cascade="all")  
  
class Address(Base): ...
```

# (results) Impedance Mismatch

What is the type of table below?

```
table = conn.execute("SELECT * FROM big_table")
```

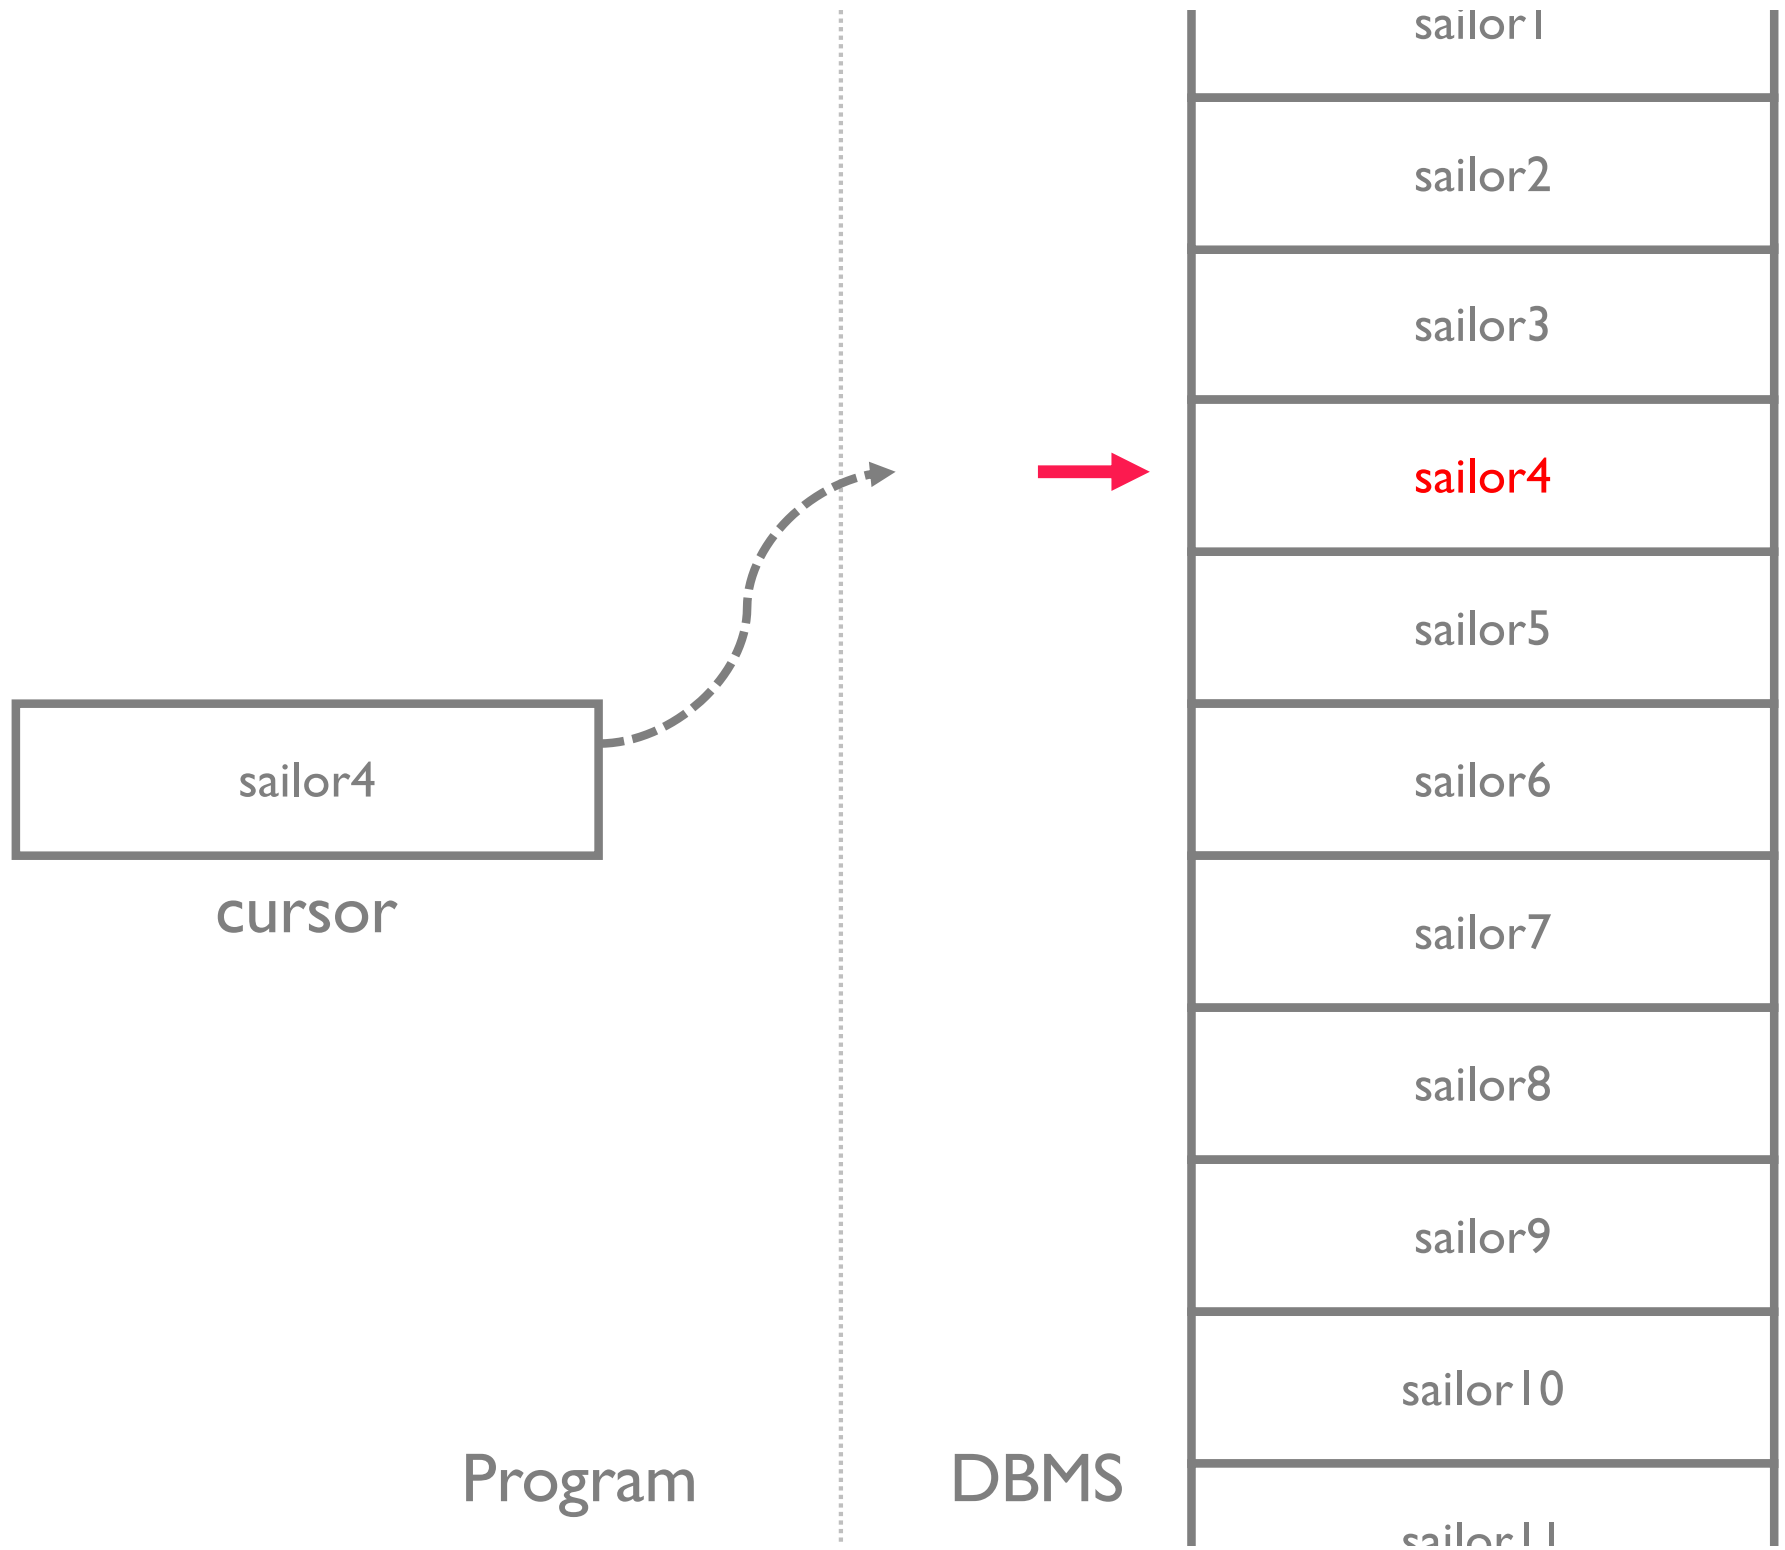
Cursor over the Result Set

similar to an iterator

Note: relations are unordered!

Cursors have no ordering guarantees

Use ORDER BY to ensure an ordering





cursor


Program



DBMS

sailor1
sailor2
sailor3
sailor4
sailor5
sailor6
sailor7
sailor8
sailor9
sailor10
sailor11





cursor

Program



sailor1
sailor2
sailor3
sailor4
sailor5
sailor6
sailor7
sailor8
sailor9
sailor10
sailor11

DBMS

# (results) Impedance Mismatch

Cursor similar to an iterator (next() calls) `cursor =`

```
cursor = conn.execute("SELECT * FROM T")
```

Core cursor attributes/methods (names may differ)

```
rowcount
```

```
attributes()
```

```
prev()
```

```
next()
```

```
get(idx)
```

# (results) Impedance Mismatch

Cursor similar to an iterator (next() calls)

```
cursor = conn.execute("SELECT * FROM T")
cursor.rowcount() # 1000000
cursor.fetchone() # (0, 'foo', ...)
for row in cursor: # iterate over the rest
    print row
```

Actual Cursor methods vary depending on implementation

# (functions) Impedance Mismatch

What about functions?

```
def add_one(val):  
    return val + 1
```

```
conn1.execute("SELECT add_one(1)")
```

Would need to embed a language runtime into DBMS

Many DBMSes support runtimes e.g., python

Can register User Defined Functions (UDFs)

# (constraints) Impedance Mismatch


DB Constraints often duplicated throughout program

JS

```
email = get_email_input();  
if (/@/.test(email))  
    error("must be a valid email");
```

Email \*

aoeu

 Must be a valid email

DBMS

```
CREATE TABLE Users (  
    email text CHECK(email ~ '@')  
)
```

# (constraints) Impedance Mismatch

ORMs let you define basic constraints


```
class Person(models.Model):  
    ...  
    first_name = models.CharField(max_length=30)  
    last_name = models.CharField(max_length=30, null=True)
```

```
CREATE TABLE myapp_person (  
    ...  
    first_name varchar(30) NOT NULL,  
    last_name  varchar(30)  
);
```

# (constraints) Impedance Mismatch

Third-party constraint libraries e.g., Great Expectations

```
expect_column_values_to_be_null(...)  
expect_column_to_exist(...)
```



Data  
sources

# Modern Database APIs

Examples: DryadLinq, SparkSQL

DBMS executor in same language (dotNET, Spark) as app code

Tricky:

- what happens to language impedance?
- what happens to exception handling?
- what happens to host language functions?

```
val lines  = spark.textFile("logfile.log")
val errors = lines.filter(_ startswith "Error")
val msgs   = errors.map(_ .split("\t")(2))

msgs.filter(_ contains "foo").count()
```



# Security

SQL Injection

Privacy vs Security

Access Controls and GRANT

Encryption

# SQL Injection

Pass *sanitized* values to the database

```
args = ('Dr Seuss', '40')
conn1.execute(
    "INSERT INTO users(name, age) VALUES(%s, %s)",
    args)
```

Pass in a tuple of query arguments

DBAPI library will *properly escape* input values

Most libraries support this

*Never construct raw SQL strings*

# SQL Injection

Why pass values using query parameters?

```
name = "eugene"
```

```
conn1.execute(  
    "SELECT * FROM users WHERE name=%s", name)
```

```
conn1.execute(  
    "SELECT * FROM users WHERE name='{name}'".format(name=name))  
  
SELECT * FROM users WHERE name='eugene'
```

# SQL Injection

Why pass values using query parameters?

```
name = "eugene';\nDELETE * FROM users;--"
```

```
conn1.execute(  
    "SELECT * FROM users WHERE name=%s", name)
```

```
conn1.execute(  
    "SELECT * FROM users WHERE name='{name}'".format(name=name))
```

```
SELECT * FROM users WHERE name='eugene';  
DELETE * FROM users;  
--'
```




Image: Daniel Rehn/Flickr, from *INPUT* mag

FEATURES | By [Joseph Cox](#) | Nov 20 2015, 8:00am

# The History of SQL Injection, the Hack That Will Never Go Away

Over 15 years after it was first publicly disclosed, SQL injection is still the number one threat to websites.

[https://motherboard.vice.com/en\\_us/article/aekzez/the-history-of-sql-injection-the-hack-that-will-never-go-away](https://motherboard.vice.com/en_us/article/aekzez/the-history-of-sql-injection-the-hack-that-will-never-go-away)

FILTER



## Running an SQL Injection Attack - Computerphile

Computerphile ✓ 1.6M views • 2 years ago

Just how bad is it if your site is vulnerable to an SQL Injection? Dr Mike Pound shows us how they work. Cookie Stealing: ...



## SQL Injection Attack Tutorial (2019)

HackHappy • 76K views • 8 months ago

SQL Injection attacks are still as common today as they were ten years ago. Today I'll discuss what are SQLi and how you can ...

CC



## SQL Injection Basics Demonstration

Imperva • 360K views • 9 years ago

Imperva presents an educational video series on Application and Database Attacks in High Definition (HD)



## Hacking Websites with SQL Injection - Computerphile

Computerphile ✓ 1.4M views • 5 years ago

Websites can still be hacked using SQL injection - Tom explains how sites written in PHP (and other languages too) can be ...

CC



## DEFCON 17: Advanced SQL Injection

Christiaan008 • 220K views • 8 years ago

Speaker: Joseph McCray Founder of Learn Security Online SQL Injection is a vulnerability that is often missed by web application ...

# SQL Injection


<http://w4111.github.io/inject>

code on github:

`w4111/w4111.github.io/src/injection/`

Use query parameters by passing form values as arguments to `execute()`  
Library sanitizes inputs automatically (and correctly!)

# SQL Injection



Project: You'll need to protect against simple SQL injections

More examples: <https://corenumb.wordpress.com/2016/05/14/mr-robot-blind-sql-injection-vulnerability/>



# Privacy vs Security

## Privacy:

- person's right to control their personal information and how it's used
- Should not release data that can be used to *infer* user's personal information
- hard to define

## Security:

- prevent unauthorized access to data
- access controls
- encryption prevents reading data even w/ access

# Access Control and GRANT

Different user accounts in w4lll DB

- staff, student, ew2493

Each user only has access privileges to read/modify subset of DB:

- Privileges: read, insert, update, delete
- Objects: Databases, Schemas, Tables, Views, Attributes
- Who? users, roles

# Access Control and GRANT

```
GRANT <privileges>  
    ON <objects>  
    TO <users/roles>
```

```
CREATE USER ew2493;  
CREATE ROLE admin;
```

```
GRANT SELECT, INSERT ON users TO admin;    // users relation  
GRANT CREATE, CONNECT ON DATABASE test TO admin;
```

```
GRANT admin to ew2493;
```

# Access Control and GRANT

How to restrict user's access to subsets of a relation or only aggregated statistics?

Combine GRANT and Views!

```
CREATE VIEW stats AS
    SELECT department, avg(salary)
    FROM costs
    GROUP BY department

GRANT SELECT ON stats TO appuser
```

# Hashing and Encryption

Prevent data from being read, even if database contents are accessed

Hashing: one way function, loses original data

- used to check equality

Data → hash() → hasheddata

hash(password) == hash(input)

Encryption: 2 way function, is reversible

- only users with key can read
- key needs to be kept safe!

Data → encrypt(key) → encdata → decrypt(key) → Data

# Hashing and Encryption

DBMS support varies

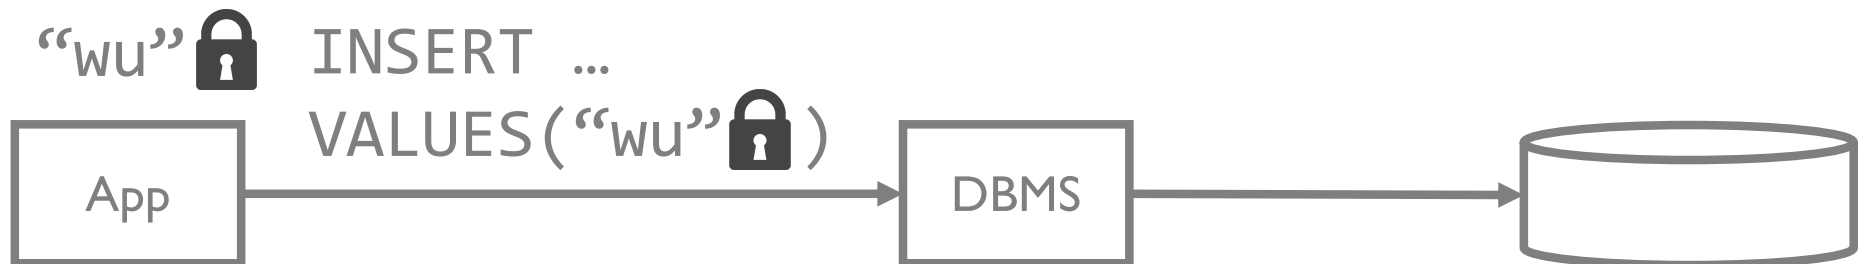
Hashing implemented as UDFs

- `INSERT INTO users VALUES(name, hash(password))`

Encryption

- encrypted hard drive, encrypted disk blocks, table, columns, ...

Application encrypts data before issuing queries



# Hashing and Encryption

DBMS support varies

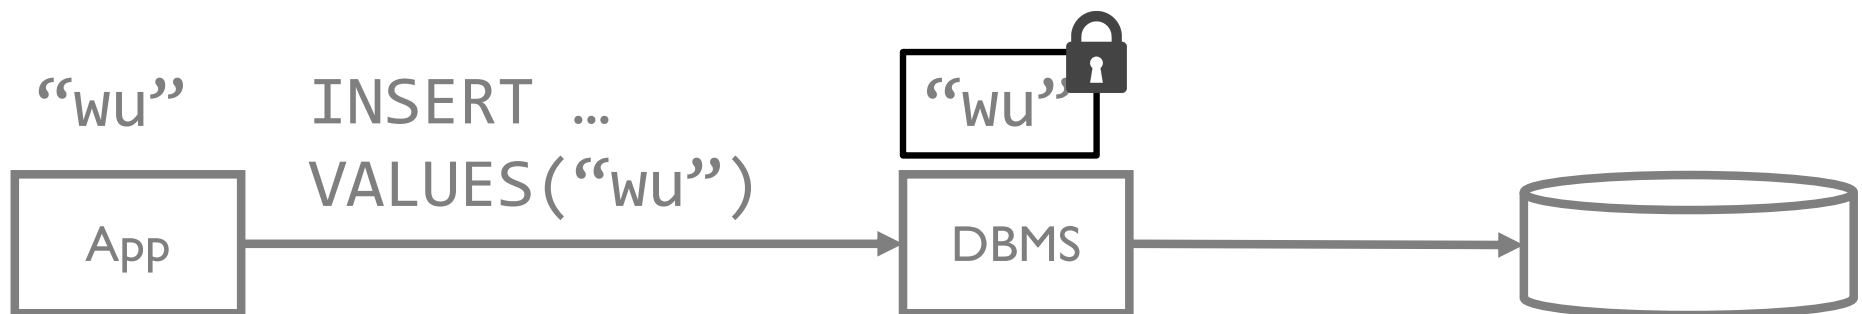
Hashing implemented as UDFs

- `INSERT INTO users VALUES(name, hash(password))`

Encryption

- encrypted hard drive, encrypted disk blocks, table, columns, ...

DBMS encrypts data received from application



# Hashing and Encryption

DBMS support varies


Hashing implemented as UDFs

- `INSERT INTO users VALUES(name, hash(password))`

Encryption

- encrypted hard drive, encrypted disk blocks, table, columns, ...

Storage encrypts everything on e.g., hard drive





# Summary

## DBAPIs

- Impedance mismatch

- Cursors

- SQL injection

Some hard queries

- More in the HW

Windows are optional material

SQL Injection: only what's in slides

# What to Understand

Why Embedded SQL is no good

Client-server vs embedded DBMSes

DBAPI components, cursors

Impedance mismatch: examples and possible solutions

SQL injection and protections

Levels of access controls and Views

Conceptual understanding of privacy & encryption