

# Data Quality

Eugene Wu

# Data Quality

## Prevention

- Normalization prevents inconsistencies
- Constraints prevent invalid values

## Detection

- Incorrect/anomalous values
- Deduplication

# Steps for a New Application

## Requirements

what are you going to build?

## Conceptual Database Design

pen-and-pencil description

## Logical Design

formal database schema

## Schema Refinement:

fix potential problems, normalization

Normalization

## Physical Database Design

use sample of queries to optimize for speed/storage

## App/Security Design

prevent security problems

# **A Relational Model of Data for Large Shared Data Banks**

**E. F. CODD**

***IBM Research Laboratory, San Jose, California***

**Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation). A prompting service which supplies such information is not a satisfactory solution. Activities of users at terminals and most application programs should remain unaffected when the internal representation of data is changed and even when some aspects of the external representation are changed. Changes in data representation will often be needed as a result of changes in query, update, and report**

# Redundancy = Bad

What was our solution?

- Break database into small relations

- Perform “good” ER modeling and SQL translation

- Kind of ... adhoc

Is there a systematic approach?

# Redundancy = Bad

<u>sid</u>	name	address	<u>hobby</u>	cost
1	Eugene	amsterdam	trucks	\$\$
1	Eugene	amsterdam	cheese	\$
2	Bob	40th	paint	\$\$\$
3	Bob	40th	cheese	\$
4	Shaq	florida	swimming	\$

people have names and addrs

hobbies have costs

people many-to-many with hobbies

What's primary key? *sid?* *sid + hobby?*

Update/insert/delete anomalies. Wastes space

# Anomalies (Inconsistencies)

## Update Anomaly

change one address, need to change all

## Insert Anomaly

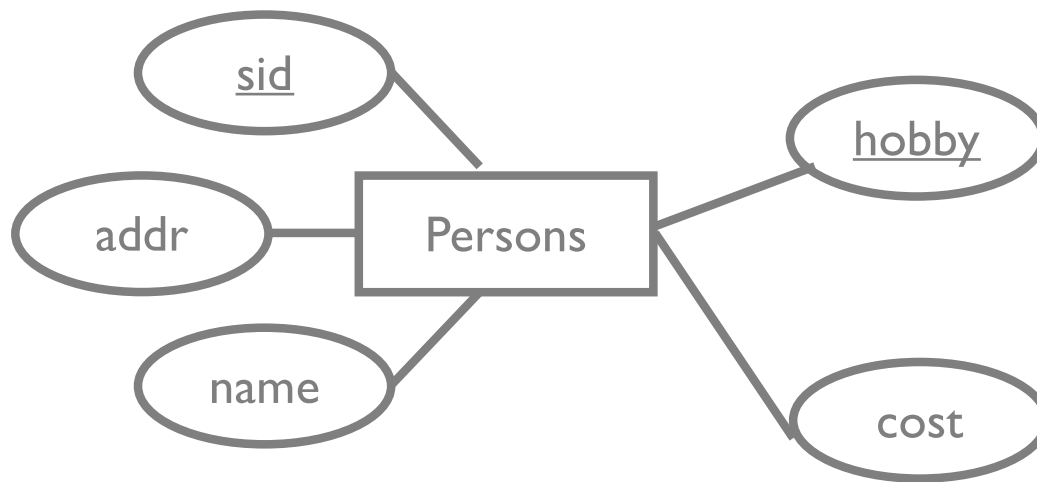
add person without hobby?  
not allowed? dummy hobby?

## Delete Anomaly

if delete a hobby. Delete the person?

**Theory Can Fix This!**

# A Possible Approach

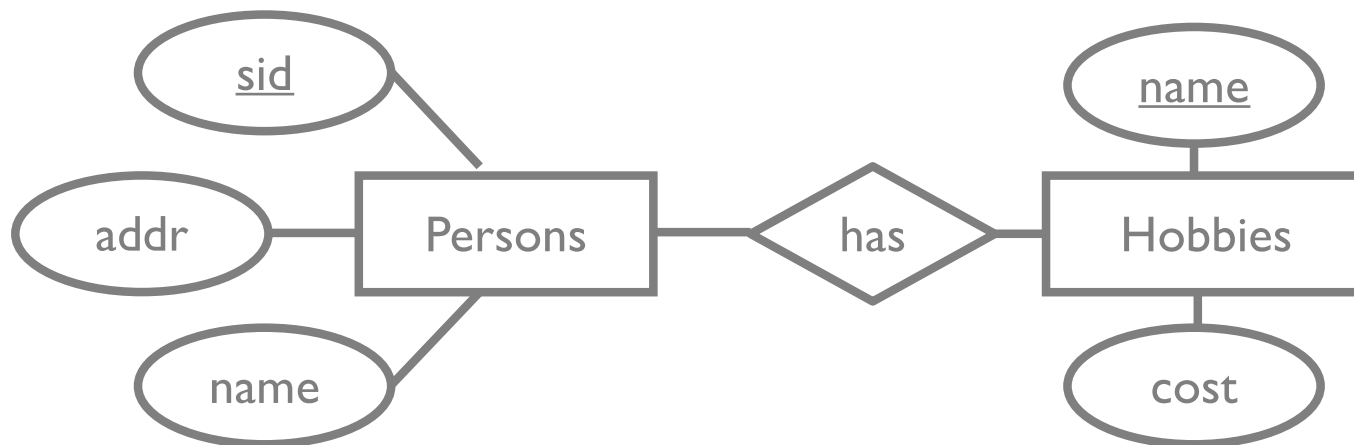


```
data(sid, addr, name, hobby, cost)
```



# A Possible Approach

ER diagram was a heuristic



We have decomposed example table into:

`person(sid, addr, name)`

`hobby(name, cost)`

`personhobby(hobbyname, sid)`

**WHY is this a good decomposition??**

# A Possible Approach

What if decompose into:

person(sid, name, addr, cost)

personhobby(sid, hobbyname)

<u>sid</u>	name	addr	cost
1	Eugene	amsterdam	\$\$
1	Eugene	amsterdam	\$
2	Bob	40th	\$\$\$
3	Bob	40th	\$
4	Shaq	florida	\$

<u>sid</u>	hobby
1	trucks
1	cheese
2	paint
3	cheese
4	swimming

but... which cost goes with which hobby?

lost information: *lossy decomposition*

# Decomposition

Replace schema R with 2+ smaller schemas that

1. each contain subset of attrs in R
2. together include all attrs in R

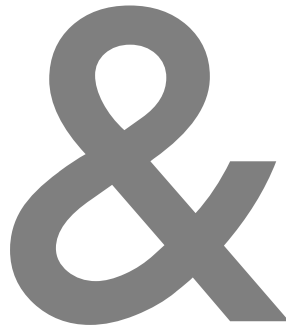
ABCD replaced with AB, BCD or AB, BC, CD

Not free – may introduce problems!

1. lossy-join: not able to recover R from smaller relations
2. non-dependency-preserving: constraints on R cannot be enforced y only looking at an individual decomposed relation
3. performance: additional joins, may affect performance

Can we systematically  
decompose our relation to

prevent  
decomposition  
problems



remove  
redundancy?

# Functional Dependencies (FD)

sid	name	address	hobby	cost
1	Eugene	amsterdam	trucks	\$\$
1	Eugene	amsterdam	cheese	\$
2	Bob	40th	paint	\$\$\$
3	Bob	40th	cheese	\$
4	Shaq	florida	swimming	\$

sid sufficient to identify name and addr, but not hobby  
e.g., exists a function  $f(\text{sid}) \rightarrow \text{name, addr}$

sid  $\rightarrow$  name, addr is a **functional dependency**

“sid determines name, addr”

“name, addr are functionally dependent on sid”

“if 2 records have the same sid, their name and addr are the same”

# Functional Dependencies (FD)

$$X \rightarrow Y$$

holds on R

if  $t_1.X = t_2.X$  then  $t_1.Y = t_2.Y$

where  $X, Y$  are subsets of attrs in R

## Examples of FDs in person-hobbies table

$\text{sid, hobby} \rightarrow \text{name, address cost}$

$\text{hobby} \rightarrow \text{cost}$

$\text{sid} \rightarrow \text{name, address}$

# Redundancy depends on FDs

consider  $R(ABC)$

no FDs: no redundancy

if  $A \rightarrow B$ : tuples with same  $A$  value means  $B$  is duplicated!



# Fun Facts

Functional Dependency is an integrity constraint  
statement about all instances of relation

Generalizes key constraints

if  $K$  is candidate key of  $R$ , then  $K \rightarrow R$

Given FDs, simple definition of redundancy

when left side of FD is not table key

Where do FDs come from?

thinking really hard aka application semantics

can't stare at database to derive (like ICs)



# Fun Facts

Functional Dependency is an integrity constraint  
statement about all instances of relation

Generalizes key constraints

if  $K$  is candidate key of  $R$ , then  $K \rightarrow R$

## **Functional Dependency Discovery: An Experimental Evaluation of Seven Algorithms**

Thorsten Papenbrock<sup>2</sup>

Jens Ehrlich<sup>1</sup>

Jannik Marten<sup>1</sup>

Tommy Neubert<sup>1</sup>

Jan-Peer Rudolph<sup>1</sup>

Martin Schönberg<sup>1</sup>

Jakob Zwiener<sup>1</sup>

Felix Naumann<sup>2</sup>

<sup>1</sup> firstname.lastname@student.hpi.uni-potsdam.de

<sup>2</sup> firstname.lastname@hpi.de

Hasso-Plattner-Institut, Prof.-Dr.-Helmert-Str. 2-3, 14482 Potsdam, Germany

# We're going to need some theory

Closure of FDs

armstrong's axioms

Principled Decomposition

BCNF

# Closure of FDs

Given

$\text{Name} \rightarrow \text{Bday}$  and  $\text{Bday} \rightarrow \text{age}$

We know

$\text{Name} \rightarrow \text{age}$

$f'$  is implied by set  $F$  if  $f'$  is true when  $F$  is true

$F^+$  **closure** of  $F$  is all FDs implied by  $F$

Can we construct this closure automatically? YES

# Closure of FDs

*Inference rules called Armstrong's Axioms*

Reflexivity      if  $Y \subseteq X$  then  $X \rightarrow Y$

Augmentation    if  $X \rightarrow Y$  then  $XZ \rightarrow YZ$  for any  $Z$

Transitivity      if  $X \rightarrow Y$  &  $Y \rightarrow Z$  then  $X \rightarrow Z$

These are **sound** and **complete** rules

sound            doesn't produce FDs not in the closure

complete        doesn't miss any FDs in the closure

# Closure of FDs

Can we compute the closure? YES. slowly  
expensive. exponential in # attributes

Can we *check* if  $X \rightarrow Y$  is in the closure of  $F$ ?

$X^+ = \text{attribute closure}$  of  $X$  (expand  $X$  using axioms)

check if  $Y$  is implied in the attribute closure

# Closure of FDs

$F = \{A \rightarrow B, B \rightarrow C, CB \rightarrow E\}$

Is  $A \rightarrow E$  in the closure?

$A \rightarrow B$

given

$A \rightarrow AB$

augmentation  $A$

$A \rightarrow BB$

apply  $A \rightarrow B$  (transitivity)

$A \rightarrow BC$

apply  $B \rightarrow C$  (transitivity)

$A \rightarrow E$

apply  $BC \rightarrow E$  (transitivity)

# We're going to need some theory

Closure of FDs

armstrong's axioms

Principled Decomposition

BCNF

# Decomposition

Eventually want to decompose  $R$  into  $R_1 \dots R_n$  wrt  $F$

We've seen issues with decomposition.

Lost Joins: Can't recover  $R$  from  $R_1 \dots R_n$

Lost dependencies

Principled way of avoiding these?



# Lossless Join Decomposition

Let's say relation R is decomposed into relations X,Y

Join the decomposed tables to get *exactly the* original

$$\pi_X(R) \bowtie \pi_Y(R) = R$$

Lossless wrt F if and only if  $F^+$  contains

$$X \cap Y \rightarrow X \text{ or } Y \cap X \rightarrow Y$$

intersection of X,Y is a key for one of them

# Lossless Join Decomposition

Lossless wrt  $F$  if and only if  $F^+$  contains

$$X \cap Y \rightarrow X \text{ or } Y \cap X \rightarrow Y$$

intersection of  $X, Y$  is a key for one of them

FDs:  $A \rightarrow C, A \rightarrow B$

A	B	C
1	2	1
5	3	4
9	2	6



A	B
1	2
5	3
9	2

B	C
2	1
3	4
2	6



A	B	C
1	2	1
5	3	4
9	2	6
1	2	6
9	2	1

Lossy!  $AB \cap BC = B$  doesn't determine anything

# Lossless Join Decomposition


Lossless wrt  $F$  if and only if  $F^+$  contains

$$X \cap Y \rightarrow X \text{ or } Y \cap X \rightarrow Y$$

intersection of  $X, Y$  is a key for one of them


FDs:  $A \rightarrow C, A \rightarrow B$

A	B	C
1	2	1
5	3	4
9	2	6



A	B
1	2
5	3
9	2

A	C
1	1
5	4
9	6



A	B	C
1	2	1
5	3	4
9	2	6

OK

# Dependency-preserving Decomposition

$F_R$  = Projection of  $F$  onto  $R$

Subset of  $F^+$  that are “valid” for  $R$

FDs  $X \rightarrow Y$  in  $F^+$  where  $X$  and  $Y$  attrs are in  $R$

If  $R$  decomposed into relations  $A$  and  $B$ .

FDs that hold on  $A, B$  are equivalent to all FDs on  $R$

$$(F_X \cup F_Y)^+ = F^+$$

Consider  $ABCD$ ,  $C$  is key,  $AB \rightarrow C, D \rightarrow A$

BCNF decomposition:  $BCD, DA$

$AB \rightarrow C$  doesn't apply to either table. Not dependency preserving.

# We're going to need some theory

Closure of FDs

armstrong's axioms

Principled Decomposition

**BCNF**

# BCNF

Relation R in BCNF has *no redundancy* wrt FDs  
(FDs are all key constraints)

F: set of functional dependencies over relation R

for  $(X \rightarrow Y)$  in  $F^+$   
Y is in X *or*  
X is a superkey of R

Is this in BCNF?

$\text{sid} \rightarrow \text{name}$

sid	hobby	name
x	y <sub>1</sub>	z
x	y <sub>2</sub>	?

# BCNF

Relation R in BCNF has *no redundancy* wrt FDs  
(FDs are all key constraints)

F: set of functional dependencies over relation R

for  $(X \rightarrow Y)$  in  $F^+$   
Y is in X *or*  
X is a superkey of R

## Functional Dependencies

$SH \rightarrow NAC$  (sid, hobby  $\rightarrow$  name, addr, cost)  
 $H \rightarrow C$   
 $S \rightarrow NA$

## What's in BCNF?

SHNAC	NO
SNA, SHC	NO
SNA, HC, SH	YES

# BCNF

Relation R in BCNF has *no redundancy* wrt FDs  
(FDs are all key constraints)

F: set of functional dependencies over relation R

for  $(X \rightarrow Y)$  in  $F^+$   
Y is in X *or*  
X is a superkey of R

Remember that  $F^+$  is *all* possible FDs valid over R!

Given  $A \rightarrow B$ ,  $B \rightarrow C$ , then  $A \rightarrow C$  is in F and should be checked



# BCNF

Suppose we have

Client, Office  $\rightarrow$  Account

Account  $\rightarrow$  Office

What's in BCNF?

R(Account, Client, Office)

R(Account, Office)   R(Client, Account)

Where did  $CO \rightarrow A$  go? *Lost a Functional Dependency*

BCNF does not preserve FDs

# BCNF Decomposition

while BCNF is violated

  R with FDs  $F_R$

  if an FD  $X \rightarrow y$  violates BCNF       $y$  is single attr

    turn R into  $R - y$  &  $Xy$

ABCDE

key A,  $BC \rightarrow A$ ,  $D \rightarrow B$ ,  $C \rightarrow D$

Consider  $D \rightarrow B$

DB is in BCNF

ACDE is in BCNF

Decompose into DB, ACDE

$D \rightarrow B$

A is key,  $C \rightarrow A$

Seems like we lost  $BC \rightarrow A$ !

# BCNF Decomposition in English

BCNF is based on checking *all* FDs (in  $F^+$ ) for a given relation.

Simple procedure:

1. For each attribute or combination of attributes in R. Call it *attrs*
2. Does the closure of *attrs* contain all attributes in R?
3. If yes: continue. Else: decompose
4. Recall that  $attrs \rightarrow attrs$  is trivially true.

Given  $A \rightarrow B$ ,  $B \rightarrow C$ , and relation ABCD

Consider A:	it determines B, and C, but not D
Decompose using $A \rightarrow B$ :	AB, ACD
Consider A for ACD:	it determines C, but not D
Decompose using $A \rightarrow C$ :	AB, AC, AD

# Other Normal Forms

Two different criterias for decomposing a relation

Boyce Codd Normal Form (BCNF)

No redundancy, may lose dependencies

3rd Normal Form (3NF)

May have redundancy, no decomposition problems

Common

1<sup>st</sup>, 2<sup>nd</sup>, 4<sup>th</sup>, ... normal forms

# Common Data Quality Issues (incomplete list)

## Value Errors

- Finding outliers
- Robustness to outliers
- Deal with outliers/value errors (Imputation)

## Entity Resolution

# Value Errors

Values in an attribute that appear wrong

- missing
- placeholder values e.g., 0, -9999, “NA”, “Empty”
- “anomalous” values (outliers)

Why are outliers bad?

- causes estimates (AVG) to be arbitrarily wrong
- $\text{AVG}(2, 2, 2, 2, 5) = 2.6$
- $\text{AVG}(2, 2, 2, 2, 100) = 21.6$
- $\text{AVG}(2, 2, 2, 2, 1000) = 201.6$

# Finding Outliers

How to detect (numeric) outliers?

- Fundamentally a modeling problem.
- What's a “normal” value? Assume a distribution

Example using age

- Assume a normal distribution
- Fit age values to distribution (compute avg, stddev)
- Low-probability ages (outside of  $\text{mean} \pm \text{std}$ ) labeled outliers

# Robustness to Outliers

Can the query be resilient to outliers?

Breakdown point:

- % of incorrect values before statistic is arbitrarily wrong
- AVG: one value
- Median: 50%

1, 2, 3, 4, 5, 6, 7, 8, 100

AVG          15.1

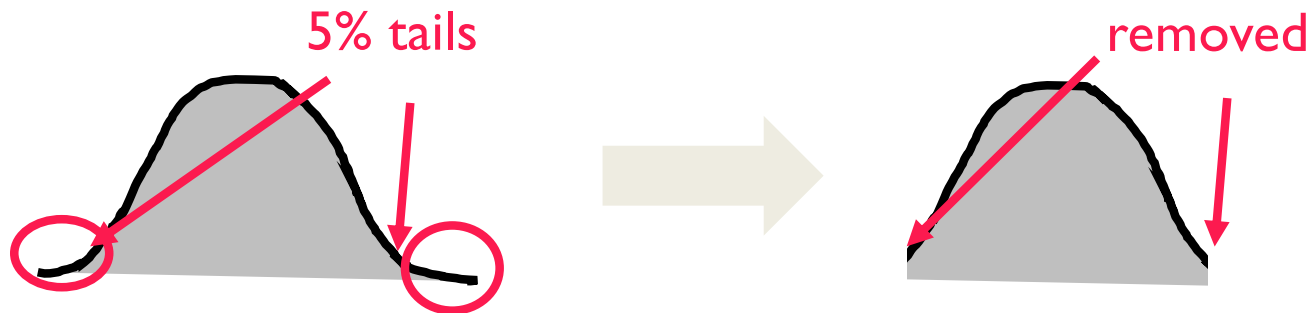
Median      5      Sort and return middle value



# Deal With Outliers

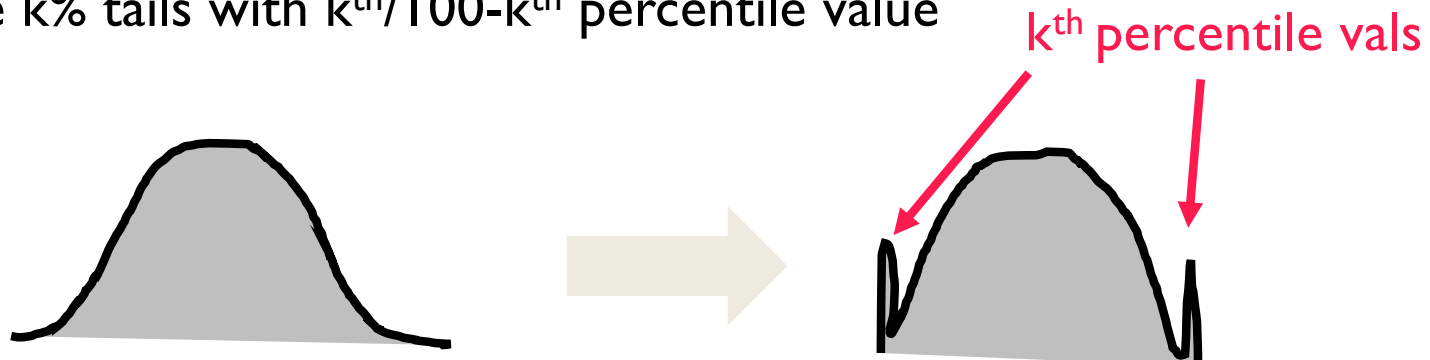
## Trimming 5%

- remove top and bottom 5% percentile (the “tails”)
- $p_5, p_{95}$



## Winsorize

- Replace extreme values with closest value
- Replace  $k\%$  tails with  $k^{\text{th}}/100$ - $k^{\text{th}}$  percentile value



# Imputing Missing Value

Fill in with “likely” values

Why?

- Some apps don't accept nulls
- Lead to bias/errors

How to define “likely”?

- default values (mean/median, user-defined val)
- interpolate (if ordered data)
- build a model and predict

# Entity Resolution

Same entity, different “physical representations”

Example: text attributes due to e.g., misspellings

- New York, NYC, New York City, NY City, ...
- Different forms of string similarity

```
SELECT *  
FROM cities c1, cities c2  
WHERE levenshtein(c1.name, c2.name) < 3
```

# Entity Resolution

Example: different rows refer to the same entity.

- Seller 1: (123, “iphone max 12Pro,” \$1000)
- Seller 2: (00123, “IPhone12 Pro MAX”, “85722.40 Rubles”)
- Fundamentally classification problem. `is_same(row1, row2) → T/F`

```
SELECT row1, row2
FROM data as d1, data as d2
WHERE is_same(d1.*, d2.*)
```

# Entity Resolution

Very hard problem! No “Solution”.

10000 records. 100,000,000 possible matches!

3 main steps

1. **Blocking**: group data into subsets (may overlap)
2. **Matching**: within each group:
  1. choose distance function between rows
  2. cluster rows using distance function.

Hope that same entities in same cluster, and  
different entities in different clusters

# Entity Resolution

Blocking:

- Extract grouping features/attributes A
- Form groups based on same/similar A values
- Each group is a “block”

Simple approach: group by an attribute

Example for strings: group by shared n-gram

2-gram for “eugene”: eu, ug, ge, en, ne

All words containing “eu” in one block

Choose n via trial-and-error

# Entity Resolution: Matching

## Distance metrics

### Single attribute

- string similarity functions
- synonym dictionary: fast = quick
- knowledge-base: *mandarin* closer to *orange* than banana
- embeddings

### Multi-attribute:

- linear combination of attribute distances  
(4111, “intro to db”), (COMS4111, “Introducton to DB”),  
(4112, “db impl”)
- train a model if you have labels

# Summary

## Normalization

- Functional dependencies & redundancy
- FD closures: Armstrong's axioms
- Proper Decomposition into BCNF

## Data Quality

- Value errors.
- Entity resolution



# Summary

## Trade-off

- Accidental redundancy is really really bad
- But adding lots of joins can hurt performance

## In practice:

- List FDs
- normalize,
- decide which decompositions to skip as needed

# What you should know

## Purpose of normalization

- Anomalies
- Decomposition problems
- Functional dependencies & axioms

## BCNF

- properties
- algorithm

## Data Quality (level of lecture)

- Value errors: How to detect, avoid, fix
- Entity resolution: What it is, why it's hard, overall steps

# Why The Term “Normalize Relations”?

Designing a database involves a process called normalization where the data model is broken-down according its smallest granularity. Invented by Ted Codd and Chris Date, the term normalization was borrowed from President Richard Nixon normalizing relations with China in the 1970's. Codd stated that if Nixon could normalize relations, then so could the relational model.