


Help with Visualization Research!

MS and Undergrad student research, took Intro to Vis in the past. Hoping to submit to the VIS conference in a few weeks.

User study to understand what people remember when reading visualizations with captions

If interested in helping, take their survey!
<https://www.surveymonkey.com/r/SD99FBR>



United States has more than 800,000 new daily COVID cases in mid-January

Administrivia

HW3 due

HW4 out: access costs and optimizer

Project 2 out: more SQL!

L8

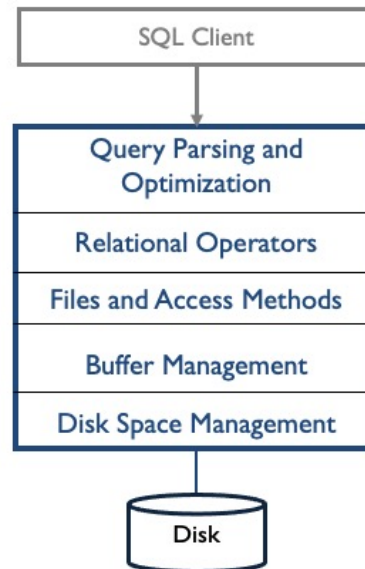
Disk, Storage, and Indexing

Eugene Wu

DBMS Overview

Each layer provides a simple abstraction to layers above it, and makes assumptions about layers below it.

Requires careful design and assumptions for performance




Classic architecture. Most large systems from Oracle, IBM etc use this. Optimized for disk.

Newer systems from new hardware may change the organization, or tweak components, however these concepts are still there.

DBMS Overview

Applications interact with
SQL Client

```
db.execute('''  
  SELECT a, b  
  FROM S, T  
  WHERE S.c = T.c''')
```



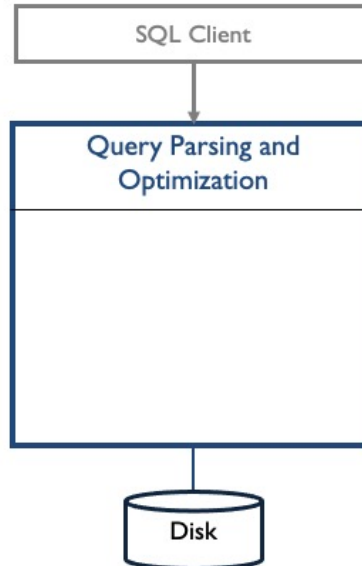

Classic architecture. Most large systems from Oracle, IBM etc use this. Optimized for disk.

Newer systems from new hardware may change the organization, or tweak components, however these concepts are still there.

DBMS Overview

Parse, check, and verify the SQL query

Turn into efficient query plan



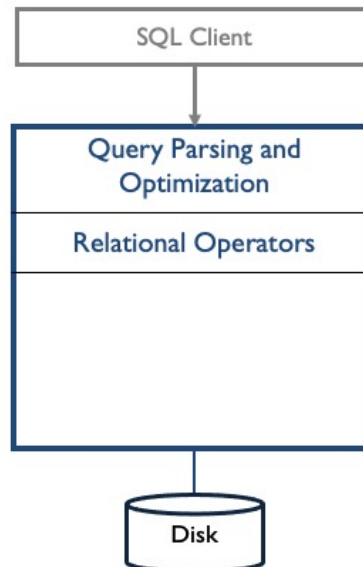

Classic architecture. Most large systems from Oracle, IBM etc use this. Optimized for disk.

Newer systems from new hardware may change the organization, or tweak components, however these concepts are still there.

DBMS Overview

A query is represented as a relational data flow

Each operator is a specific implementation



Classic architecture. Most large systems from Oracle, IBM etc use this. Optimized for disk.

Newer systems from new hardware may change the organization, or tweak components, however these concepts are still there.

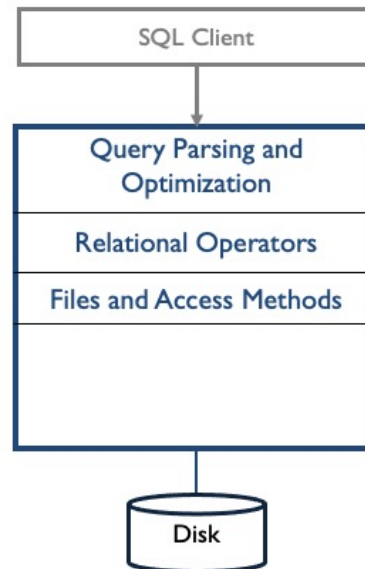
What is a heap scan reading from? What about index scan?

DBMS Overview

Organizes tables, indexes, records as groups of pages in a “logical file”

API:

- Operators ask for records
- Logical files help read and write bytes on pages



Classic architecture. Most large systems from Oracle, IBM etc use this. Optimized for disk.

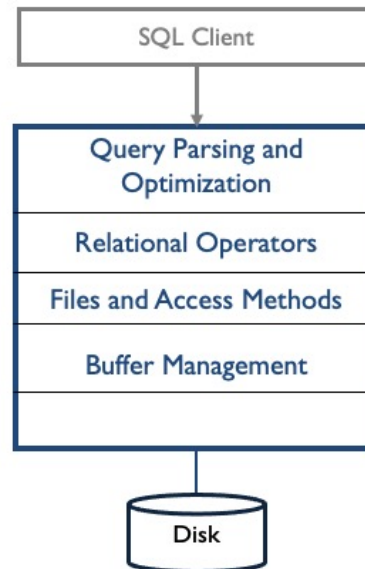
Newer systems from new hardware may change the organization, or tweak components, however these concepts are still there.

DBMS Overview

Not all pages can fit into RAM.

Buffer manager provides illusion that all pages are accessible.

Files simply ask for pages.




Classic architecture. Most large systems from Oracle, IBM etc use this. Optimized for disk.

Newer systems from new hardware may change the organization, or tweak components, however these concepts are still there.

DBMS Overview

Physically read and write bytes on one or more storage devices (hard drives, SSDs, etc)

Storage performance properties dictate the design of layers above.




Classic architecture. Most large systems from Oracle, IBM etc use this. Optimized for disk.

Newer systems from new hardware may change the organization, or tweak components, however these concepts are still there.

DBMS Overview

Layers help manage engineering complexity.

Requires assumptions about performance of lower layers. (cost models)



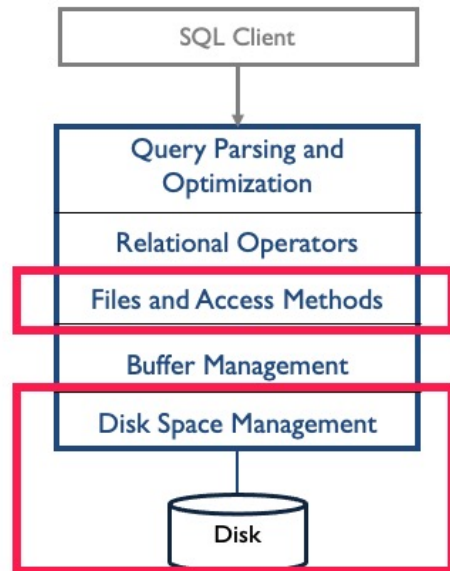
Classic architecture. Most large systems from Oracle, IBM etc use this. Optimized for disk.

Newer systems from new hardware may change the organization, or tweak components, however these concepts are still there.

DBMS Overview

Physically read and write bytes on one or more storage devices (hard drives, SSDs, etc)

Storage performance properties dictate the design of layers above.



Classic architecture. Most large systems from Oracle, IBM etc use this. Optimized for disk.

Newer systems from new hardware may change the organization, or tweak components, however these concepts are still there.

Storage Devices

Hard Drives vs SSDs vs RAM

Locality (random vs sequential)

Reads vs Writes

Performance and Monetary Costs

Why not store all in RAM? \$\$\$

Too much \$\$\$

High-end DBs: ~Petabyte (1000TB).

SQL Hyperscale: 100TB+ (2018)

Disks are ~60% cost of a production system

Main memory not persistent

Obviously important if DB stops/crashes

main-memory DBMSes discussed in advanced DB course

in many cases, you don't have petabytes of data, and main memory or SSDs are practical.

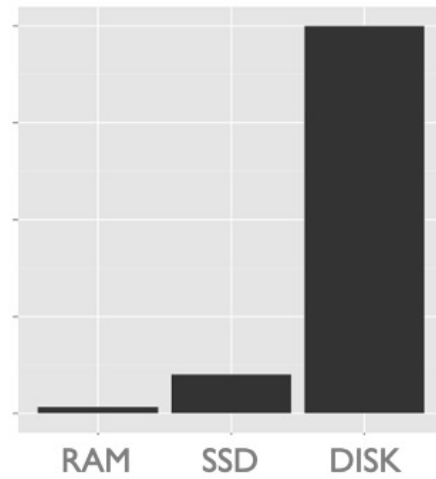
\$ Matters

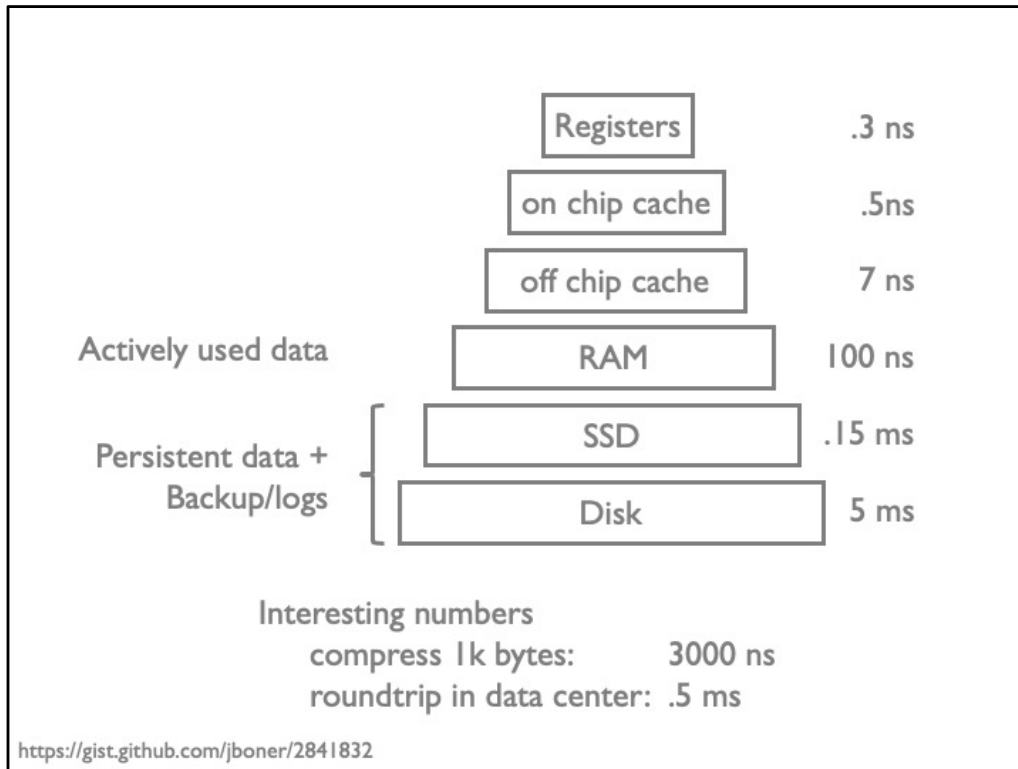
Newegg enterprise \$1000

RAM: 0.08TB

SSD: ~2TB (25x)

Disk: ~40TB (500x)





We'll focus on RAM and Disk and algorithms between the two.

It turns out what really matters is the performance ratio between the two there are some algorithms specialized to how a disk works, but for most part the types of techniques DBs use between RAM and disk can be applied in for example chip cache and RAM, and indeed many techniques such as pre-fetching are commonly used in the Os as well


L2: 14x L1

Ram: 20x L2

compression: 0.003 milliseconds

Roundtrip in a data center: 0.5 ms: 10x faster than disk seek

disk: 50k times slower than RAM



Philly = sacramento

jim gray basically wrote the book on transaction processing, the ideas of transactions, ACID, data cube, 5 minute rule

5 minute rule: The 5-minute random rule: cache randomly accessed disk pages that are re-used every 5 minutes or less.

In 2000, Gray and Shenoy applied a similar calculation for [web page caching](#) and concluded that a browser should "cache web pages if there is any chance they will be re-referenced within their lifetime."^[8]

go read his wikipedia page


[https://en.m.wikipedia.org/wiki/Jim_Gray_\(computer_scientist\)](https://en.m.wikipedia.org/wiki/Jim_Gray_(computer_scientist))

Spin speed: ~7200-15K RPM. 10K RPM ~ 6ms/rotation

Sector: multiples of pages.

More sectors further from center

Terminology: block = page in this class



disk is a stack of these platters that are spinning just like a dj turn table – just 100x faster

the platters are coated with magnetic material that is used to flip bits between 1 and 0

the data is laid out in tracks – concentric circles like trees or music record

the track is split into tiny sectors or blocks of roughly 64kb, varies by manufacturer
think of a block like a page – similar to an OS page, usually OS pages are multiple of disk blocks for nice properties

when want to read/write, you'll hear a little whirling sound, as the arm moves to position the head,

no random access, no pointers, no objects.

what's changed, has been the magnetic material on the surface of the platters, and encodings, etc, but the main thing, the physics, has not changed.

that's the only mechanical device in your computer!

API means need to

- READ: transfer page of data to disk from ram
- write: transfer page from disk to ram

Kinda slow. really slow

IS this the right api?

Time to access (read or write) a disk block

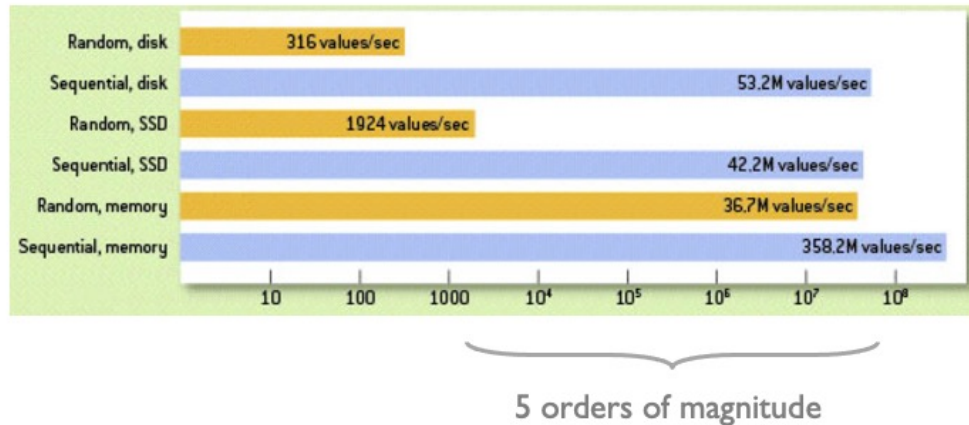
seek time	2-4 msec avg
rotational delay	2-4 msec
transfer time	0.3 msec/64kb page

Throughput

read	~150 MB/sec
write	~50 MB/sec

Key: reduce seek and rotational delays
HW & SW approaches

of 4 byte values read per second



throughput is comparable between disk and SSD! The main difference is random access and latency

SSDs

NAND memory

Small reads: 4-8k


Big writes: 1-2MB

2-3k writes before failure

wear leveling: distribute writes around

Write amplification: changing 4 bytes writes a 1-2MB chunk!

Need to think about wear, garbage collection, writing in bulk



<https://www.anandtech.com/show/2738/5>


SSDs

Reads fast

- single read time: 0.03ms
- random reads: 500MB/s
- sequential reads: 525MB/s

Writes less predictable

- single write time: 0.03ms
- random writes: 120MB/s
- sequential writes: 480MB/s



What's Best? Depends on Application

Small databases: SSD/RAM

All global daily weather since 1929: 20GB

2000 US Census: 200GB

2009 english wikipedia: 14GB

Easily fits on an SSD or in RAM

Very Big databases: Disk

Sensors easily generate TBs of data/day

Boeing 787 generates ½ TB per flight

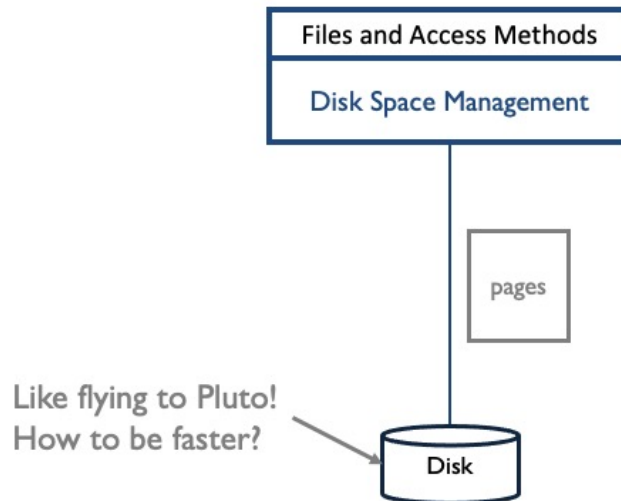
Disk has best cost-capacity ratio

SSDs help reduce read variance

when would you have cloud scale databases

if in sciences, or in practice – small number of machines, or a big desktop makes sense

Work from the bottom up



All of this is very complicated – and we DONT want to deal with sectors, or tracks, or platters.

So the abstraction use to communicate with the disk is in pages. We say we want to write or read a set of pages, and the disk controller will help manage that request.

Strategies for Fast Data Access

(think about going to the store)

Big difference between random & sequential access

- Optimize for sequential accesses

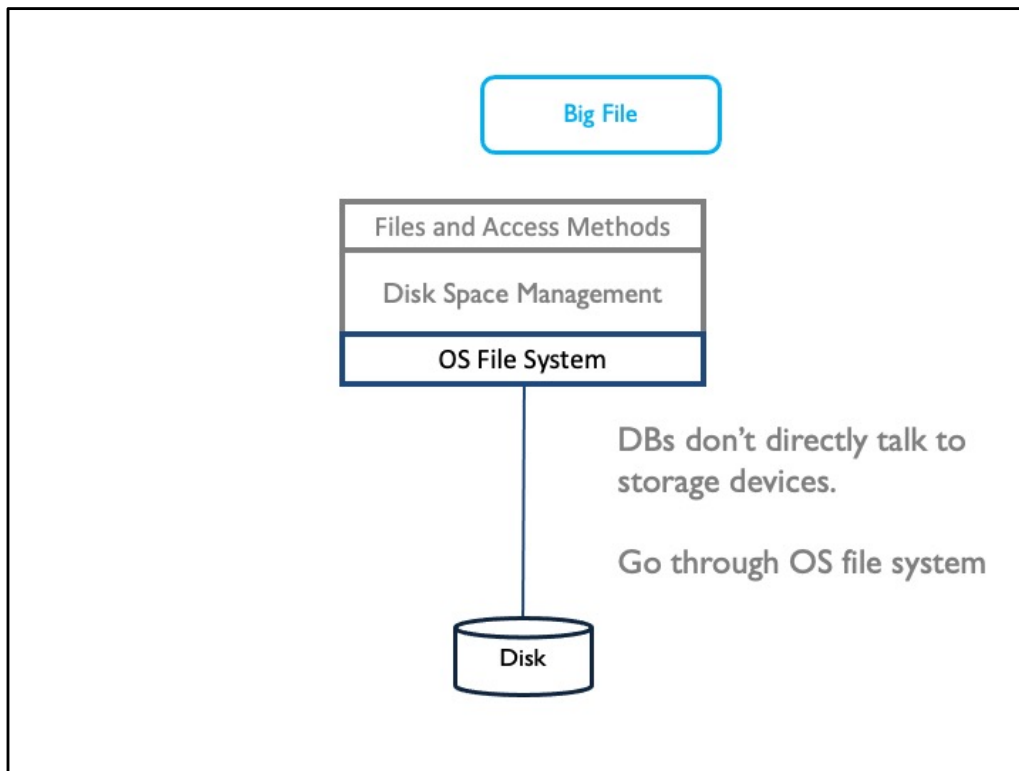
Amortize sequentially read & write big chunks of bytes

Cache popular blocks

Pre-fetch what you will need later

API

- read/write page
- read/write sequential pages
- notion of “next” page (upper layers can assume next is faster)



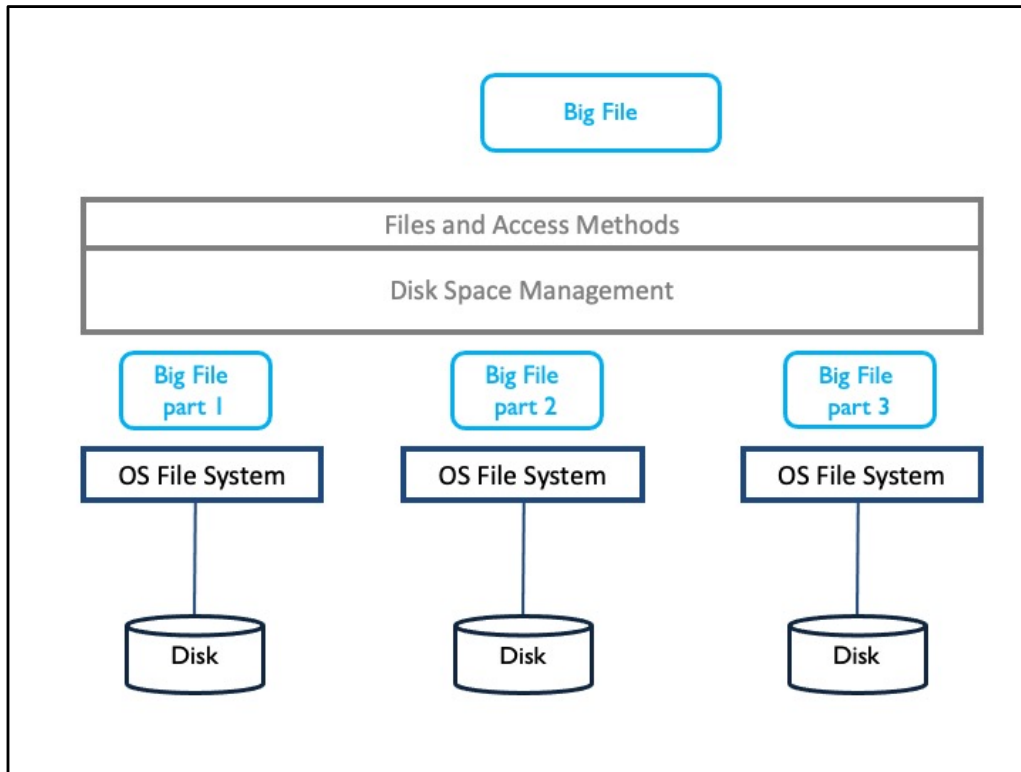
could imagine directly managing the hardware, but then you need to talk to different physical devices and deal with drivers etc.

A huge amount of operating system code is for dealing with and providing drivers for a wide range of hardware devices, so best let OS manage that and give us a file abstraction

usually, we allocate a huge amount of space on disk – usually allocated sequentially, and once we have that, use file API to read write blocks, with the understanding that the file is on disk

Higher level don't have guarantees that things will be sequential, BUT if we know things are sequential we can use better algorithms

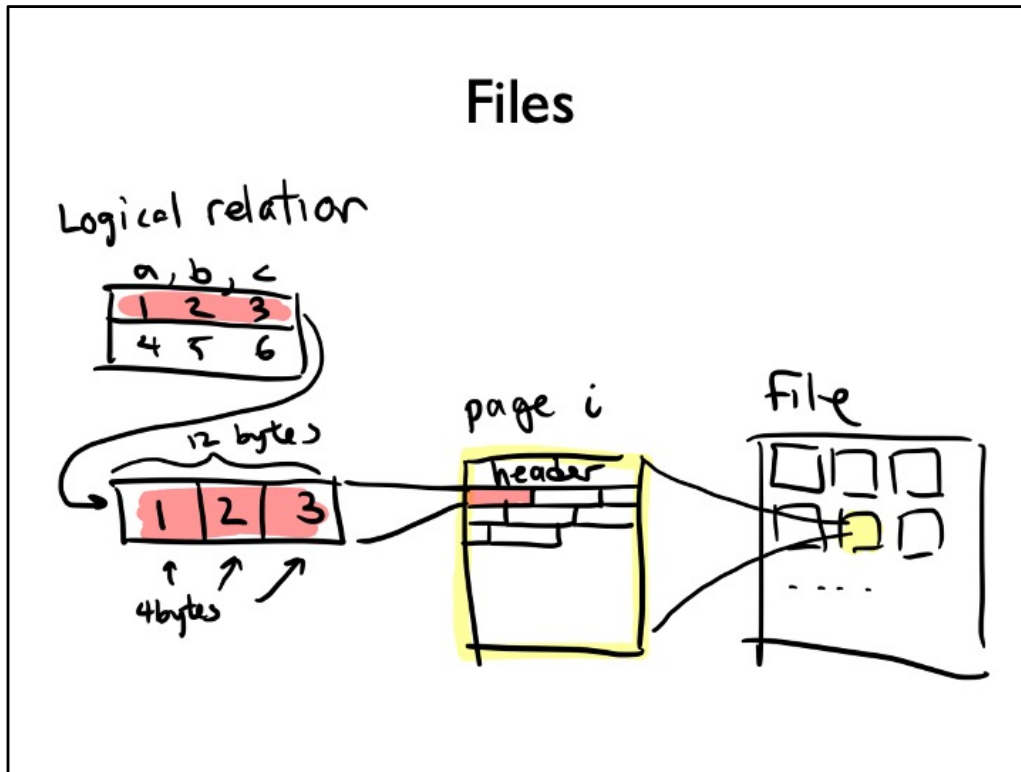
our operations will be at the level of pages



This abstraction is nice, because we are not tied to a single file system. Can partition data across multiple file systems as well, and none of our files and access methods and higher layers in the DBMS need to care.

PostgreSQL database stored in files

```
osulldyn-160-39-248-197 ~/d/p/base> cd 33883/
osulldyn-160-39-248-197 ~/d/p/b/33883> ls
112      12526    2601      2612      2659      2696      2840_vn    3456_fsm    3602
113      12538    2601_fsm  2612_fsm  2660      2699      2841      3456_vn    3602_fsm
1247     12529    2601_vn   2612_vn   2661      2701      2995      3466_vn    3602_vn
1247_fsm 12529_fsm 2602      2613      2662      2702      2995_vn    3466_vn    3603
1247_vn   12529_vn 2602_fsm  2613_vn   2663      2703      2996      3467      3603_fsm
1249      12531    2602_vn   2615      2664      2704      3079      3468      3603_vn
1249_fsm  12533    2603      2615_fsm  2665      2753      3079_fsm  3501      3604
1249_vn   12534    2603_fsm  2615_vn   2666      2753_fsm  3079_vn   3501_vn    3605
12504     12536    2603_vn   2616      2667      2753_vn   3080      3502      3606
12504_fsm 12538    2604      2616_fsm  2668      2754      3081      3503      3607
12504_vn   1255     2604_vn   2616_vn   2669      2755      3085      3534      3608
12506     1255_fsm 2605      2617      2670      2756      3118      3541      3609
12508     1255_vn   2605_fsm  2617_fsm  2673      2757      3118_vn   3541_fsm   3712
12509     1259     2605_vn   2617_vn   2674      2830      3119      3541_vn    3764
12509_fsm  1259_fsm  2606      2618      2675      2830_vn   3164      3542      3764_fsm
12509_vn   1259_vn   2606_fsm  2618_fsm  2676      2831      3256      3574      3764_vn
12511     1417     2606_vn   2618_vn   2679      2832      3256_vn   3575      3766
12513     1417_vn   2607      2619      2680      2832_vn   3257      3576      3767
12514     1418     2607_fsm  2619_fsm  2681      2833      3258      3576_vn    548
12514_fsm  1418_vn   2607_vn   2619_vn   2682      2834      3394      3596      549
12514_vn   174      2608      2620      2683      2834_vn   3394_fsm  3596_vn    826
12516     175      2608_fsm  2620_vn   2684      2835      3394_vn   3597      826_vn
12518     2187     2608_vn   2650      2685      2836      3395      3598      827
12519     2328     2609      2651      2686      2836_vn   34002     3598_vn    828
12519_fsm  2328_vn   2609_fsm  2652      2687      2837      34004     3599      PG_VERSION
12519_vn   2336     2609_vn   2653      2688      2838      34004_fsm 3600      pg_fillnode.nap
12521     2336_vn   2610      2654      2689      2838_fsm  34004_vn  3600_fsm   pg_internal.init
12523     2337     2610_fsm  2655      2690      2838_vn   34008     3600_vn
12524     2600     2610_vn   2656      2691      2839      34031     3601
12524_fsm  2600_fsm  2611      2657      2692      2840      3455      3601_fsm
12524_vn   2600_vn   2611_vn   2658      2693      2840_fsm  3456      3601_vn
osulldyn-160-39-248-197 ~/d/p/b/33883>
```



Think File == Table

need way of mapping records to pages to files

abstraction is pages, and we read and write pages

note that it's a COLLECTION. no ordering

no assumptions of WHERE the pages live, we don't care.

contrast with unix file API

- stream of bytes
- there's an ordering
- DB File is unordered

We'll have different types of files, with different organizations that make certain types of record access patterns faster or slower

Fancier files provide additional access methods for e.g., looking up records by value rather than record id

Files

Higher layers want to talk in terms of records, and files of records

File: collection of pages

Minimum API:

- insert/delete/modify record

- lookup record_id

- scan all records

Page: collection of records

- typically *fixed page sizes* (8 or 64kb in PostgreSQL)

These are logical.

Different page organizations in a file have different access costs

Units that we'll care about

Ignore CPU cost

Ignore RAM cost

B # data pages on disk for relation

R # records per data page

D avg time to read/write data page to/from disk

Simplifies life when computing costs

OK to not be exactly correct

we'll talk about non-data pages that are part of the index

ultimately this will all be important for talking about performance tradeoffs of different ways to physically represent a file, so we need some performance modeling. ignoring a lot of details including seek times, etc could always add that in

Given the above, how long does it take to read the entire relation?

How many records are in the relation?

Unordered Heap Files

Collection of records (no order)

As we add records, pages allocated

As we remove records, pages removed

To support record level ops, need to track:

- pages in file

- free space on pages

- records on page

Ok, let's design that

Super Naïve Design



Big array
of bytes

Heap is a big array of bytes.
First split into array of pages

header page (directory) with two doubly linked lists, of full pages, and not full pages
location of header page stored in a database catalog (somewhere special)

what's a pointer on the disk? pointer? no. sector of the track etc? Nope. OS will
give us a block number (disk block ID)

what's bad about this? what's this good for?

which pages have how much free space? We don't know. Need to walk through free
space linked list

how to find records? will need to scan all of the pages unless we know something
more.

Super Naïve Design



Array of
pages

Insert

4

Each each page from disk sequentially
Check if there's an open spot
Insert 4 into open spot.

header page (directory) with two doubly linked lists, of full pages, and not full pages
location of header page stored in a database catalog (somewhere special)

what's a pointer on the disk? pointer? no. sector of the track etc? Nope. OS will
give us a block number (disk block ID)

what's bad about this? what's this good for?

which pages have how much free space? We don't know. Need to walk through free
space linked list

how to find records? will need to scan all of the pages unless we know something
more.

Super Naïve Design



Array of
pages

Some problems include

- Slow to know what pages are empty, full, have space
- Slow to find a particular value
- Fragmentation

header page (directory) with two doubly linked lists, of full pages, and not full pages
location of header page stored in a database catalog (somewhere special)

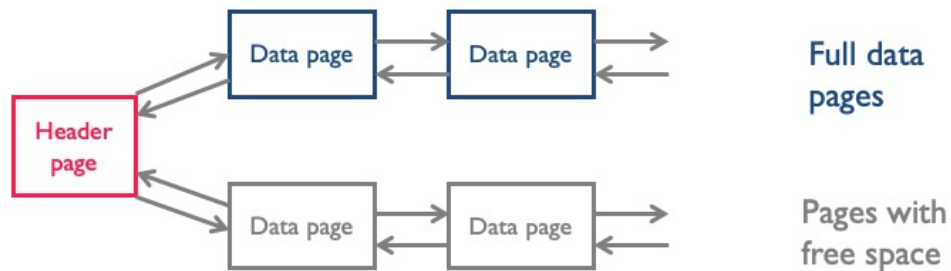
what's a pointer on the disk? pointer? no. sector of the track etc? Nope. OS will
give us a block number (disk block ID)

what's bad about this? what's this good for?

which pages have how much free space? We don't know. Need to walk through free
space linked list

how to find records? will need to scan all of the pages unless we know something
more.

Heap File



Header page info stored in catalog

Data page contains: 2 pointers, free space, data

Need to scan pages to answer any query

Which page has enough free space for 100 bytes?

header page (directory) with two doubly linked lists, of full pages, and not full pages
location of header page stored in a database catalog (somewhere special)


what's a pointer on the disk? pointer? no. sector of the track etc? Nope. OS will give us a block number (disk block ID)

what's bad about this? what's this good for?

which pages have how much free space? We don't know. Need to walk through free space linked list

how to find records? will need to scan all of the pages unless we know something more.

Use a directory




Directory entries track # free bytes on data pages

Directory is collection of pages

header pages connected together, a linked list of pointers.
each entry of directory has a pointer to a data page and how much free data, so scan directory instead of data

lots of pointers in a header page, so should be pretty small
usually good enough if using this approach



Indexes

"If I had 8 hours to
chop down a tree,
I'd spend 6 sharpening
my ax."

Abraham Lincoln

Indexes

Heap files answer any query via a sequential scan, but...

Queries use *qualifications* (predicates)

find students where class = "CS"

find students with age > 10

Indexes: file structures for value-based queries

B+-tree index (~1970s)

Hash index

Overview! Details in 4112

How would we find a record by rid? scan through the linked list until we find it.
Expectation is $\frac{1}{2}$ of all data pages

Keep in mind, indexes are designed to make things faster – with tradeoffs about what types of accesses they speed up.

It is common to use up more space to build indices than for the actual data.

In all of this, we'll be setting up to be able to compare the query costs of using each type of access method

Indexes

Defined with respect to a *search key*
don't confuse with candidate keys!!


Faster access for WHERE clauses w/ search key

```
CREATE INDEX idx1 ON users USING btree (sid)
CREATE INDEX idx2 ON users USING hash  (sid)
CREATE INDEX idx3 ON users USING btree (age,name)
```

You will play around with indexes in HW4

<https://www.postgresql.org/docs/11/sql-createindex.html>

Primary vs Secondary Index Files



Primary

Data entries contain actual tuples

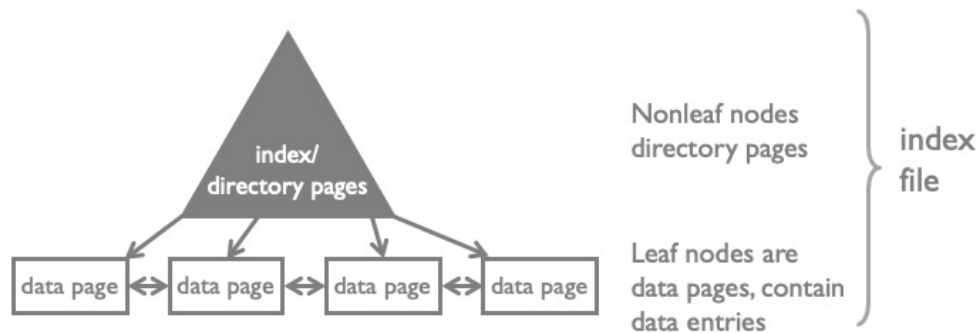
Pros: directly access data.

Secondary

Data entries contain <search key val, rid> pointers into another file

Pros: index is more compact

B+ Tree Index: disk-optimized index



Node = Page

Supports equality and range queries on search key

Self balancing (path to any leaf is almost the same)

Leaf nodes are connected via doubly linked list

Height is with respect to directory pages (the gray part of the triangle)

We saw that the directory for the heap file can reduce the cost of certain operations.

- What if we allowed multiple levels of directories?
- And kept them in sorted order on the values?

In contrast to traditional binary search trees, where each node is a single value, B+ tree nodes are pages that contain multiple values. This serves to increase the throughput when reading tree nodes.

In fact, many “key value” stores like mongo and berkeley DB are persistent B+Tree data structures

Workhorse of most DBMSes

Consists of an index structure for directing the search algorithm along with data entries as the leaf nodes that contain the actual data (same data pages as in unordered heap files)

This entire structure is composed of pages – index heap file


Terminology:

height: is wrt the index entires. So a tree of height 1 includes a single root node and data entries

B+ Tree on (age)

Non-leaf Directory pages
m index entries
m+1 pointers

Leaf Data pages
data entries/tuples



At each level in the tree, index & data page contents are sorted by search key

Query: `SELECT * WHERE age= 14`



Example of a b+ tree indexed on val (val is the *index key*)

Each non-leaf node is like a directory:

sorted list of values.

index entry = age value to compare against when searching the tree

the left and right side of each non-null index entry are pointers to the child nodes

this means there are N keys and N+1 pointers in a directory page

Here, I only show the age value of the tuples stored in the data pages


Unlike a binary tree, interior nodes are only a directory – don't store data. data is all in the leaves

How to search? Load index page and do binary search. Can do it since loaded in memory.

Index Only Queries: B+ Tree on (age)

Non-leaf Directory pages
m index entries
m+1 pointers

Leaf Data pages
data entries/tuples



At each level in the tree, index & data page contents are sorted by search key


Query: **SELECT age WHERE age > 14**
(index only!)



Note that if the data entries are <age value, rid>, then a query that projects the search key can be *index only*


B+ Tree on (age)

Note: 50 not a
data entry




Query: **SELECT * WHERE age = 55**

directory page



Unlike a binary tree, nonleaves are only a directory – don't store data. data is all in the leaves

B+ Tree on (age)



When directory page is full, can't add more pointers
Split directory page and reorganize tree
(don't need to know details of how splits are done)

If we add more data, let's say we have 2 additional pages of data, then this directory page is full and we can't add more pointers, and so we need to split it up in order to index the new pages


B+ Tree on (age)



When directory page is full, can't add more pointers
Split directory page and reorganize tree
(don't need to know details of how splits are done)

We might split it up to index the pages this way, but this is not a tree
notice that the index entry 50 has disappeared. This is because the index entries
don't contain data.

B+ Tree on (age)



**When directory page is full, can't add more pointers
Split directory page and reorganize tree
(don't need to know details of how splits are done)**


The details of how we do this don't matter, but it's self balancing, so any combination of inserts, updates, and deletes end up with a balanced tree (meaning the left and right children are roughly the same amount of data)

typically hundred(s) of items in a page

What's a benefit of the doubly linked list at the bottom of the b+tree?

It supports range queries as well. Here we go to the page with 20 (or smallest number larger than 20) and scan along the leaf pages

B+ Tree on (age)



Query: `SELECT * WHERE age > 20`


The details of how we do this don't matter, but it's self balancing, so any combination of inserts, updates, and deletes end up with a balanced tree (meaning the left and right children are roughly the same amount of data)

typically hundred(s) of items in a page

What's a benefit of the doubly linked list at the bottom of the b+tree?

It supports range queries as well. Here we go to the page with 20 (or smallest number larger than 20) and scan along the leaf pages

Composite Keys: B+ Tree on (age, name)



How do the following queries use the index on (age, name)?

```
SELECT age WHERE age = 14
SELECT *   WHERE age < 18 AND name < 'monica'
SELECT age WHERE name = 'bobby'
```

Composite key

Note that (17, alice) < (17, monica) even though both have 17.

Q1: index only, use first part of the index key

Terminology




Keep free space for fast inserts
Reduce "fill factor" parameter

Directory Page



Fanout ("branching factor")



Example using 8kb pages

How many entries in the index's data pages?

fill-factor: ~66%

~300 entries per directory page

height 2: $300^3 \sim 27$ Million entries assuming 300 tuples/pg

height 3: $300^4 \sim 8.1$ Billion entries assuming 300 tuples/pg

Top levels often in memory

height 2 only 300 pages ~2.4MB

height 3 only 90k pages ~750MB

Cool B+Tree viz: <https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>

8 kb pages, integer entries and integer pointers (8 + 8 bytes) = 500 entries in a directory

60% fill factor is 300 entries

Recall that the DBMS allocates a big chunk of memory to use for itself.

Given some standard caching policy, what are the chances the root node will be in memory?

Well it's accessed on every lookup, so it's likely in memory

What about the next level? Doesn't take much space, and probably accessed frequently, so also in memory.

So for a 27M entry table, only 1 IO to access the data page. Pretty good!

Logarithmic data structures are good when the constants are large (e.g., fanout)

Height = length of path from root to leaf

when height = 2, why do we use 300^3 ? recall height is just the index entry pages

level 3 == root (level 1) -> level 2 -> level 3 -> data entry

For Real?

problem. We have run TB range databases on Raspberry PI with 4GB of RAM and hammered that in benchmarks (far exceeding the memory capacity). The interesting thing here is that the B+Tree nature means that the upper tiers of the tree were already in memory, so we mostly ended up with a single page fault per request.

<https://ayende.com/blog/196161-c/re-are-you-sure-you-want-to-use-mmap-in-your-database-management-system>

Hash Index on age

Hash function

$$h(v) = v \% 3$$

Hash buckets
containing
data pages

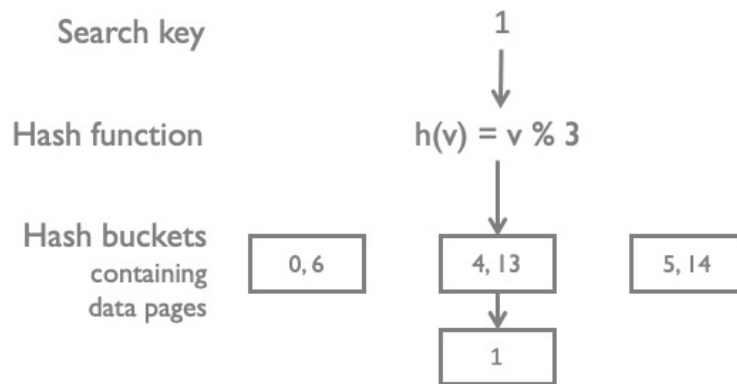
0, 6

4, 13

5, 14

hash indexes: single or multiple hash functions. Array of data pages. compute hash using hash function to get data page and insert into it. If data page is full, add an overflow page

INSERT Hash Index on age



INSERT Hash Index on age


Search key 11

Hash function $h(v) = v \% 3$

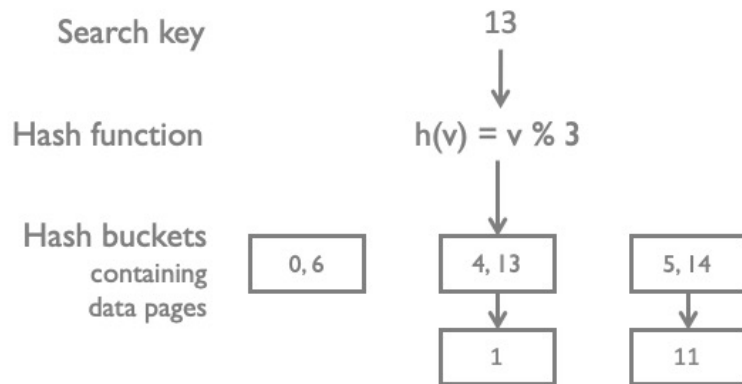
Hash buckets
containing
data pages



INSERT Hash Index on age



SEARCH Hash Index on age



Good for equality selections

Index = data pages + overflow data pages

Hash function $h(v)$ takes as input the *search key*

In terms of metadata, how is this different than B+ trees?

What types of queries is this good for compared to B+ trees?

Costs

Three file types

Heap, B+ Tree, Hash

Indexes can be primary or secondary

Operations we care about

Scan all data `SELECT * FROM R`

Equality `SELECT * FROM R WHERE x = I`

Range `SELECT * FROM R WHERE x > 10 and x < 50`

Insert tuple

Delete tuple `DELETE WHERE ...`

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything				
Equality				
Range				
Insert				
Delete				

B # data pages
D time to read/write page
M # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD			
Equality	0.5BD			
Range	BD			
Insert	2D			
Delete	Search + D			

Heap File

equality on a key. How many results?

- B** # data pages
- D** time to read/write page
- M** # pages in range query

B: total number of data pages in table

M: if doing a range query, we are fetching M pages

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD		
Equality	0.5BD	$D(\log_2 B)$		
Range	BD	$D(\log_2 B + M)$		
Insert	2D	Search + BD		
Delete	Search + D	Search + BD		

Heap File

equality on a key. How many results?

Sorted File

files compacted after deletion

B # data pages

D time to read/write page

M # pages in range query

we assume that the heap is sorted on the query predicate attribute, otherwise it's as good as an unordered heap

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD	1.25BD	
Equality	0.5BD	$D(\log_2 B)$	$D(\log_{80} B + 1)$	
Range	BD	$D(\log_2 B + M)$	$D(\log_{80} B + M)$	
Insert	2D	Search + BD	$D(\log_{80} B + 2)$	
Delete	Search + D	Search + BD	$D(\log_{80} B + 2)$	

Heap File

equality on a key. How many results?

Sorted File

files compacted after deletion

B+ Tree

100 entries/directory page

80% fill factor

B # data pages

D time to read/write page

M # pages in range query

why does scanning take 1.2BD? (see the assumptions for B+Tree in the slide)

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD	1.25BD	1.25BD
Equality	0.5BD	$D(\log_2 B)$	$D(\log_{80} B + 1)$	D
Range	BD	$D(\log_2 B + M)$	$D(\log_{80} B + M)$	1.25BD
Insert	2D	Search + BD	$D(\log_{80} B + 2)$	2D
Delete	Search + D	Search + BD	$D(\log_{80} B + 2)$	2D

Heap File

equality on a key. How many results?

Sorted File

files compacted after deletion

B+ Tree

100 entries/directory page

80% fill factor

Hash index

80% fill factor when computing scans

assumes no overflow when computing equality/insert/delete.

B # data pages

D time to read/write page

M # pages in range query

can you even perform range query with a hash index?

why is hash 1.2BD for range query? don't know the exact domain of the values!

$1 < x < 10$

try 1, 2, 3, 4, ... 10?

what about 1.5?

Where do B, D, M come from?

Estimated from more basic info

Assuming (we will be clear about this)

- fanout = number of directory entries
- pointer in secondary index same size as directory entry

Given:

- p: page size
- r: record size
- d: directory entry size
- f: fill factor
- n: # records

Estimate for primary and secondary index:

- size
- height
- access cost

Primary B+ Index

Page Size	$p = 100$	records/page	$p / r = 10$
Record Size	$r = 10$	direntries/page	$p / d = 20$
Dir entry Size	$d = 5$	fanout:	20
Fill Factor	$f = 100\%$	# data pages	$n/(p/r) = 800$
# Records	$n = 8000$	height	$\log_{20} 800 = 3$

Cost to look up a single record is
3 for directory pages + 1 data page

Secondary B+ Index

Page Size	$p = 100$	records/page	$p / r = 10$
Record Size	$r = 10$	direntries/page	$p / d = 20$
Dir entry Size	$d = 5$	fanout:	20
Fill Factor	$f = 100\%$	# data pages	$n/(p/d) = 400$
# Records	$n = 8000$	height	2

Cost to look up a single record is
2 for directory pages + 1 data page + 1 pointer lookup

How to pick?

Depends on your queries (workload)

Which relations?

Which attributes?

Which types of predicates ($=$, $<$, $>$)

Selectivity

Insert/delete/update queries? how many?

selectivity

- why wouldn't you use hash index for a range query?
- what is equality but selectivity?
- 0.1 selectivity = will return ~10% of the tuples

if all of your queries are inserts, then a heap file may make the most sense

if all of your queries are primary key equality accesses, then hash table may be a good idea

How to choose indexes?

Considerations

- which relations should have indexes?
- on what attributes?
- how many indexes?
- what type of index (hash/tree)?

called Physical database design problem

which relations are we accessing? Are they already fast? Or are they slow and an index would help?

(amount of improvement to queries on relation)

attributes: recall that b+tree or hash depend on the search key

Composite search key or single attribute search key?

Naïve Algorithm

```
get query workload
group queries by type
for each query type in order of importance
    calculate best cost using current indexes
    if new index IDX will further reduce cost
        create IDX
```

Why not create every index?

```
update queries slowed down (upkeep costs)
takes up space
```

workload

in many databases, the index sizes can often be much much larger than the actual data, so that queries go faster.

What if you don't use update queries?

High level guidelines

Check the WHERE clauses

attributes in WHERE are search/index keys

equality predicate → hash or tree index

range predicate → tree index

Multi-attribute search keys supported

order of attributes matters for range queries

may enable queries that don't look at data pages (*index-only*)

didn't talk about index-only

Summary

Design depends on economics, access cost ratios

Disk still dominant wrt cost/capacity ratio

Many physical layouts for files

- same APIs, difference performance

- remember physical independence

Indexes

- Structures to speed up read queries

- Multiple indexes possible

- Decision depends on workload

Things to Know

How a hard drive works and its major performance characteristics

The storage hierarchy & differences between RAM, SSD, Hard drives

What files, pages, and records are, and how different than UNIX model

Heap File data structure

B+ tree and Hash indexes

Performance characteristics of different file organizations

Given statistics, figure out directory size, index height, access cost