

Last Lecture

Eugene Wu

Administrivia

Review session 5/6 2-4PM 480 Mudd

- Thank the staff!

Project 2 due 5/9 10AM. *no late days!*

Extra credit opportunities

- project 2, recursive queries
- write a tutorial
- draw a concept/topic

Grading

Normalize all assignments to [0-1]

Weighed sum based on assignment weights

Compute cut-offs, ~ B+ avg

Then add in extra credit

Agenda

1. Project Winners
2. Modern Databases
3. Q&A / Review

DBMSes in the Wild

Classic Relational

\$\$: Oracle, IBM, Microsoft, Teradata, EMC, etc

Free: MySQL, PostgreSQL

New Relational

In-Memory, Column-store, Streaming

Non-traditional

Search (Google, Bing, Lucene), Scientific, Geographic

NoSQL

Big Data: Hadoop, Spark, etc

Key-value: Mongo, BerkeleyDB, Cassandra, etc

DBMS-as-a-Service

Microsoft Azure, Amazon Redshift/RDS, etc...

DBMSes in the Wild

Classic Relational

\$\$: Oracle, IBM, Microsoft, Teradata, EMC, etc

Free: MySQL, PostgreSQL

New Relational

In-Memory, Column-store, Streaming

Non-traditional

Search (Google, Bing, Lucene), Scientific, Geographic

NoSQL

Big Data: Hadoop, Spark, etc

Key-value: Mongo, BerkeleyDB, Cassandra, etc

DBMS-as-a-Service

Microsoft Azure, Amazon Redshift/RDS, etc...

Modern Database Systems

Hardware changes affect

Compression is good → Column stores


Large scale aggregation queries

“Data Warehouses”





Memory is cheap → In-memory stores

Transactional systems




One Size Fits All



One Size **Does Not** Fits All



One Size **Does Not** Fits All



Data Warehouses

Store all historical data for future analysis


Sales by month over past 20 years

Clicks by youth in texas

Cost by product component

Defacto standard at any company

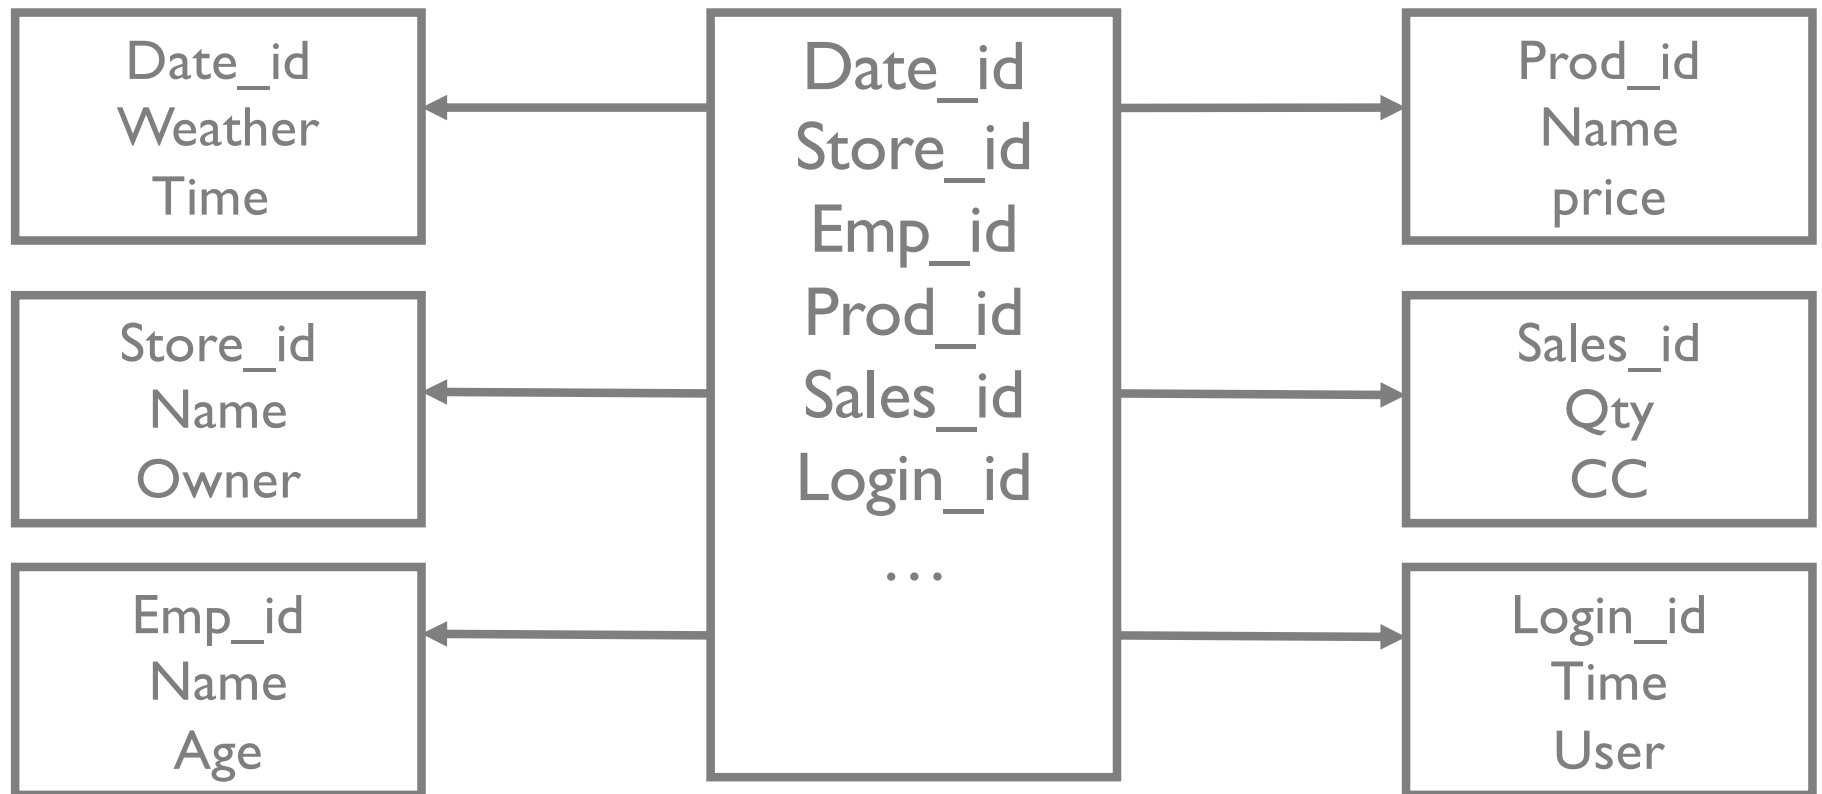
Star Schema



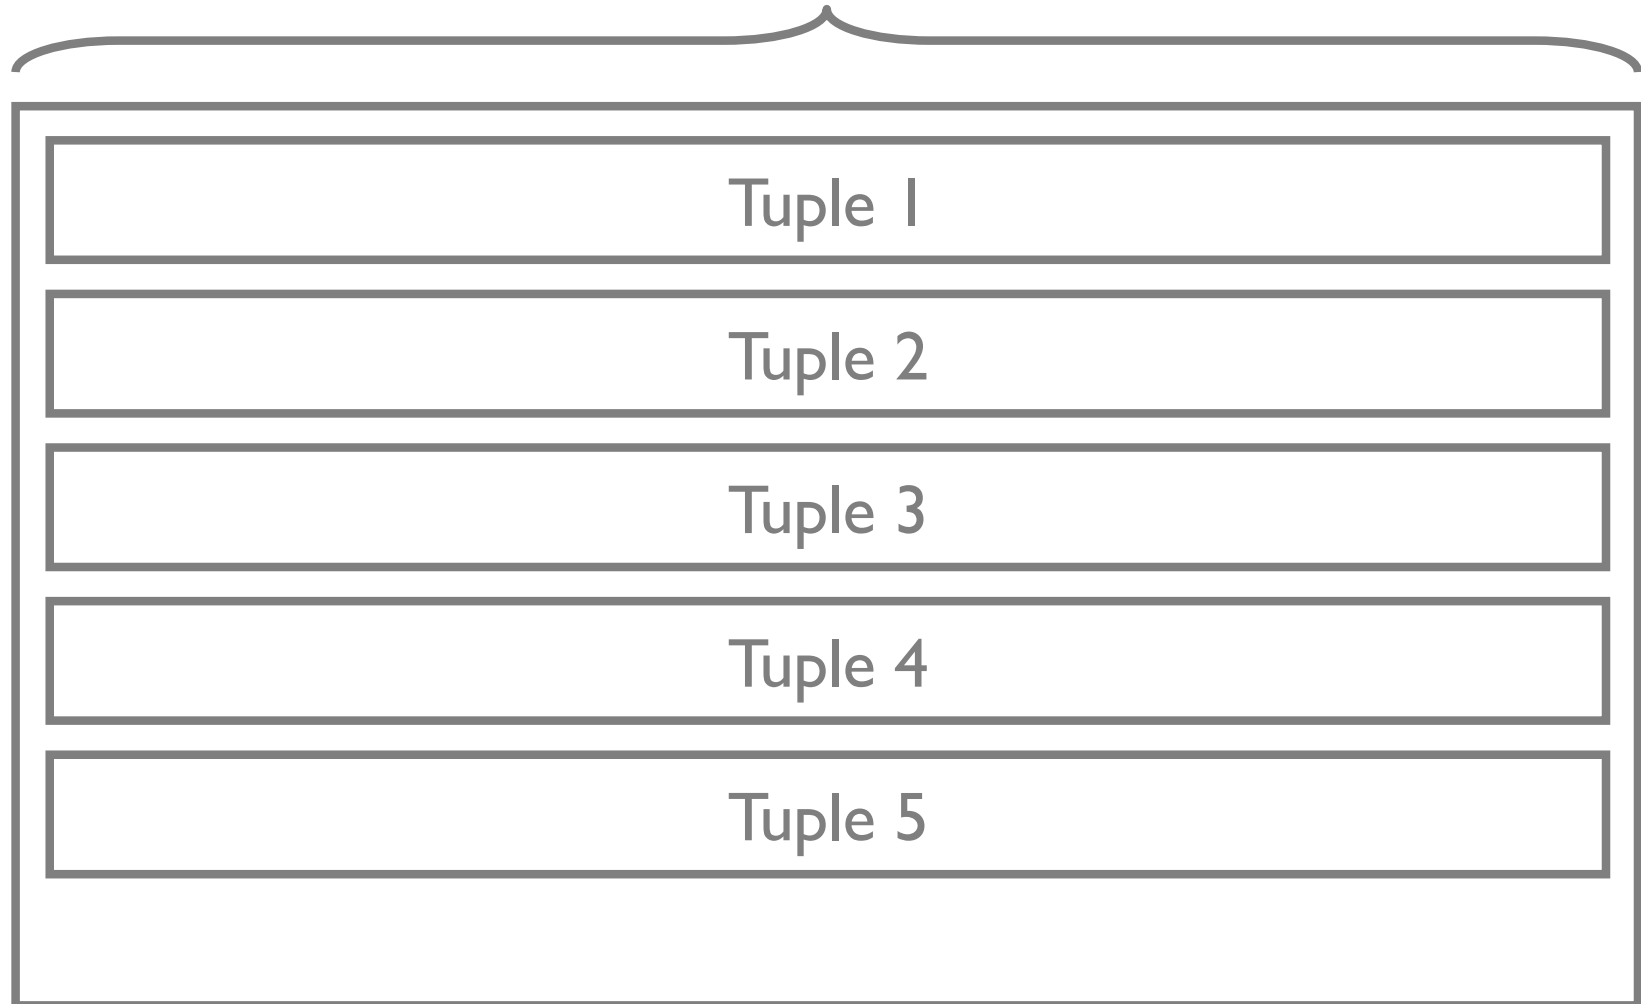
Star Schema


Fact table is “fat”

Queries access ~6 attrs




100 attributes






16x less data read. Unfair advantage.

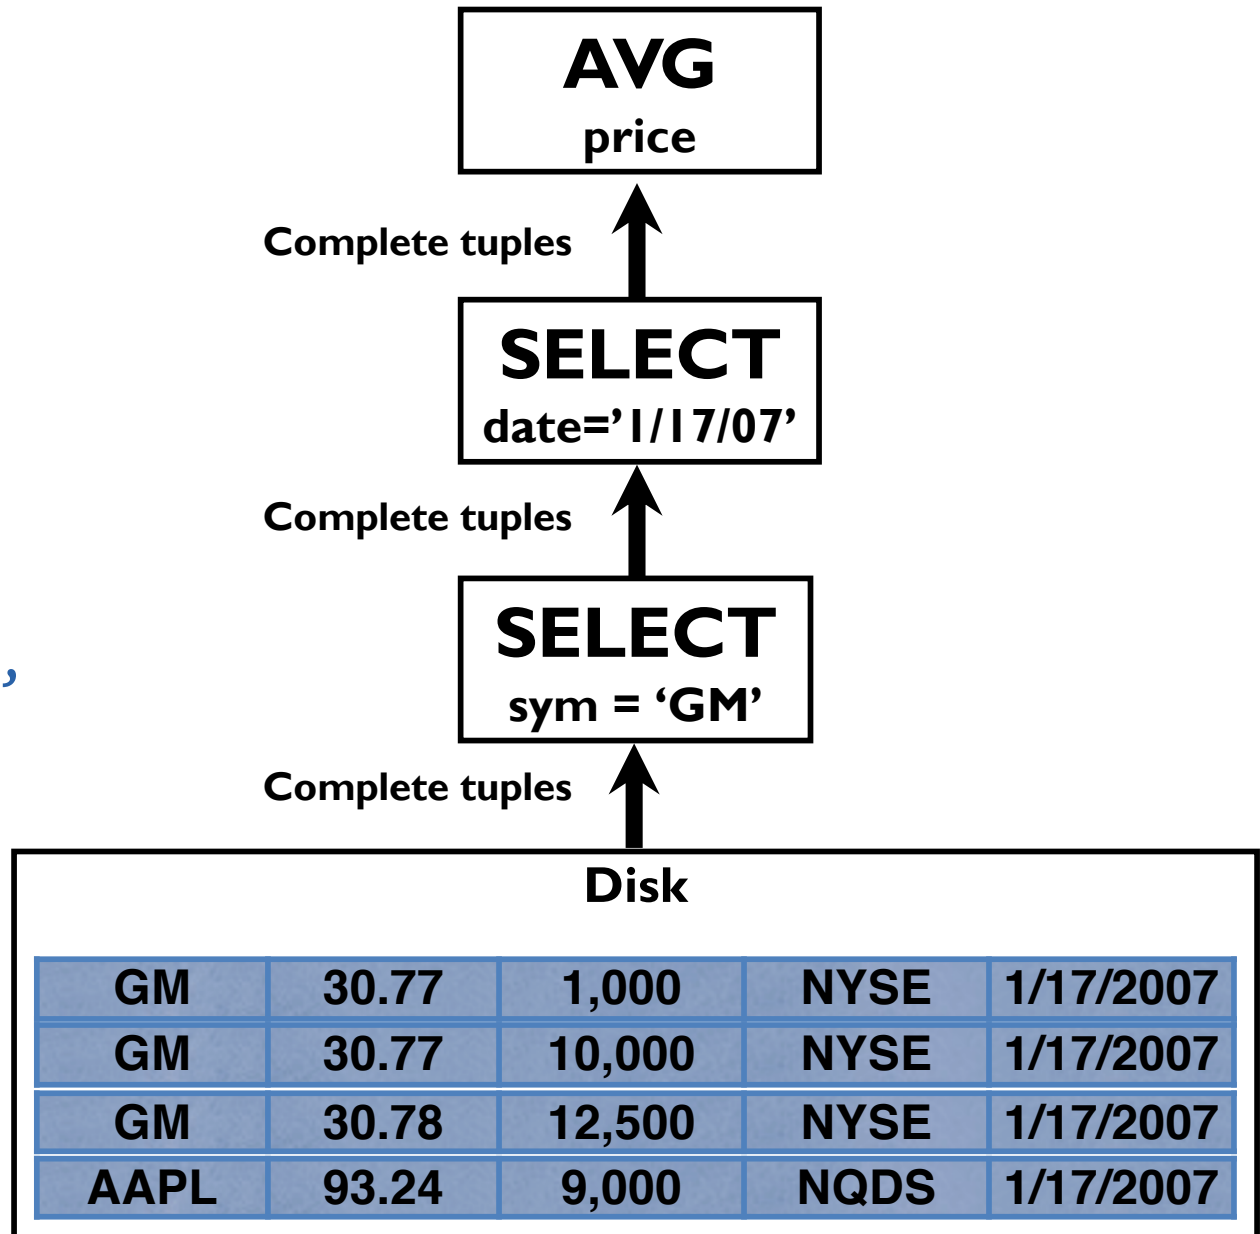


16x less data read. Unfair advantage.
Compression better on single column
Execute on compressed data



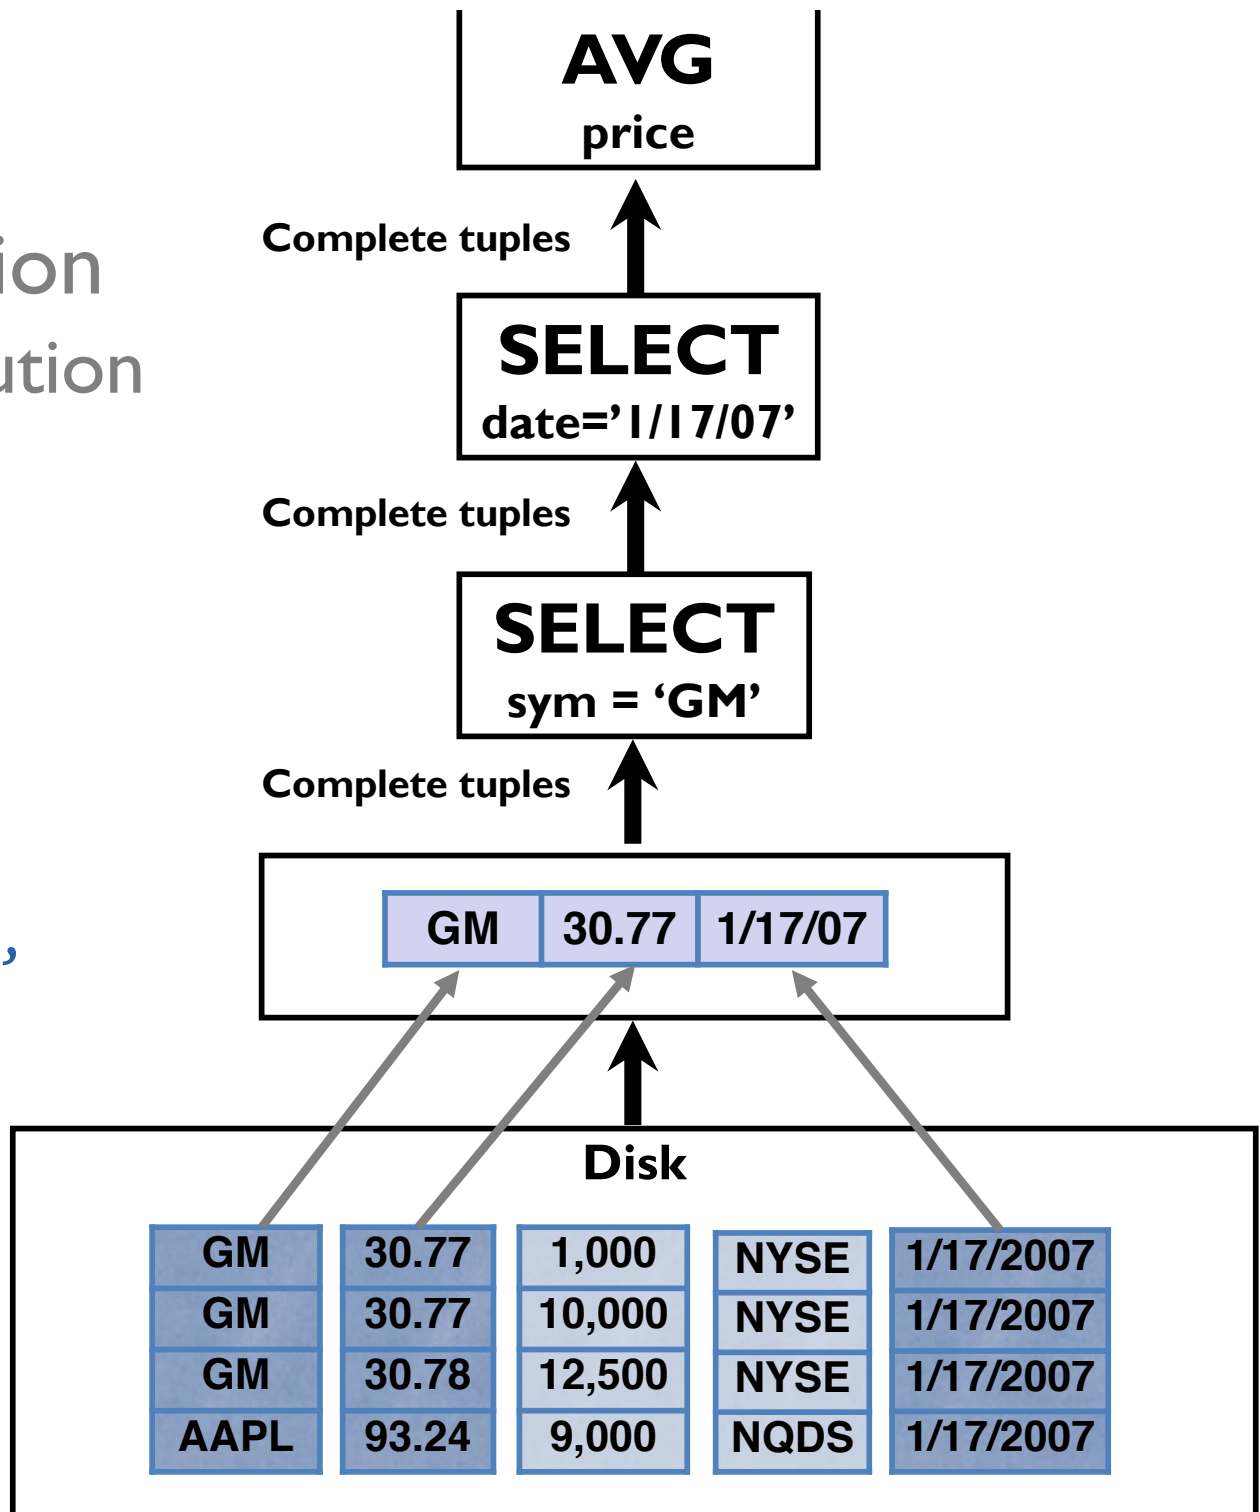
Traditional DBMS

```
SELECT avg(price)
FROM tickstore
WHERE symbol = 'GM'
AND date = '1/17/2007'
```



Naïve:
Early Materialization
Row oriented execution

```
SELECT avg(price)
FROM tickstore
WHERE symbol = 'GM'
AND date = '1/17/2007'
```

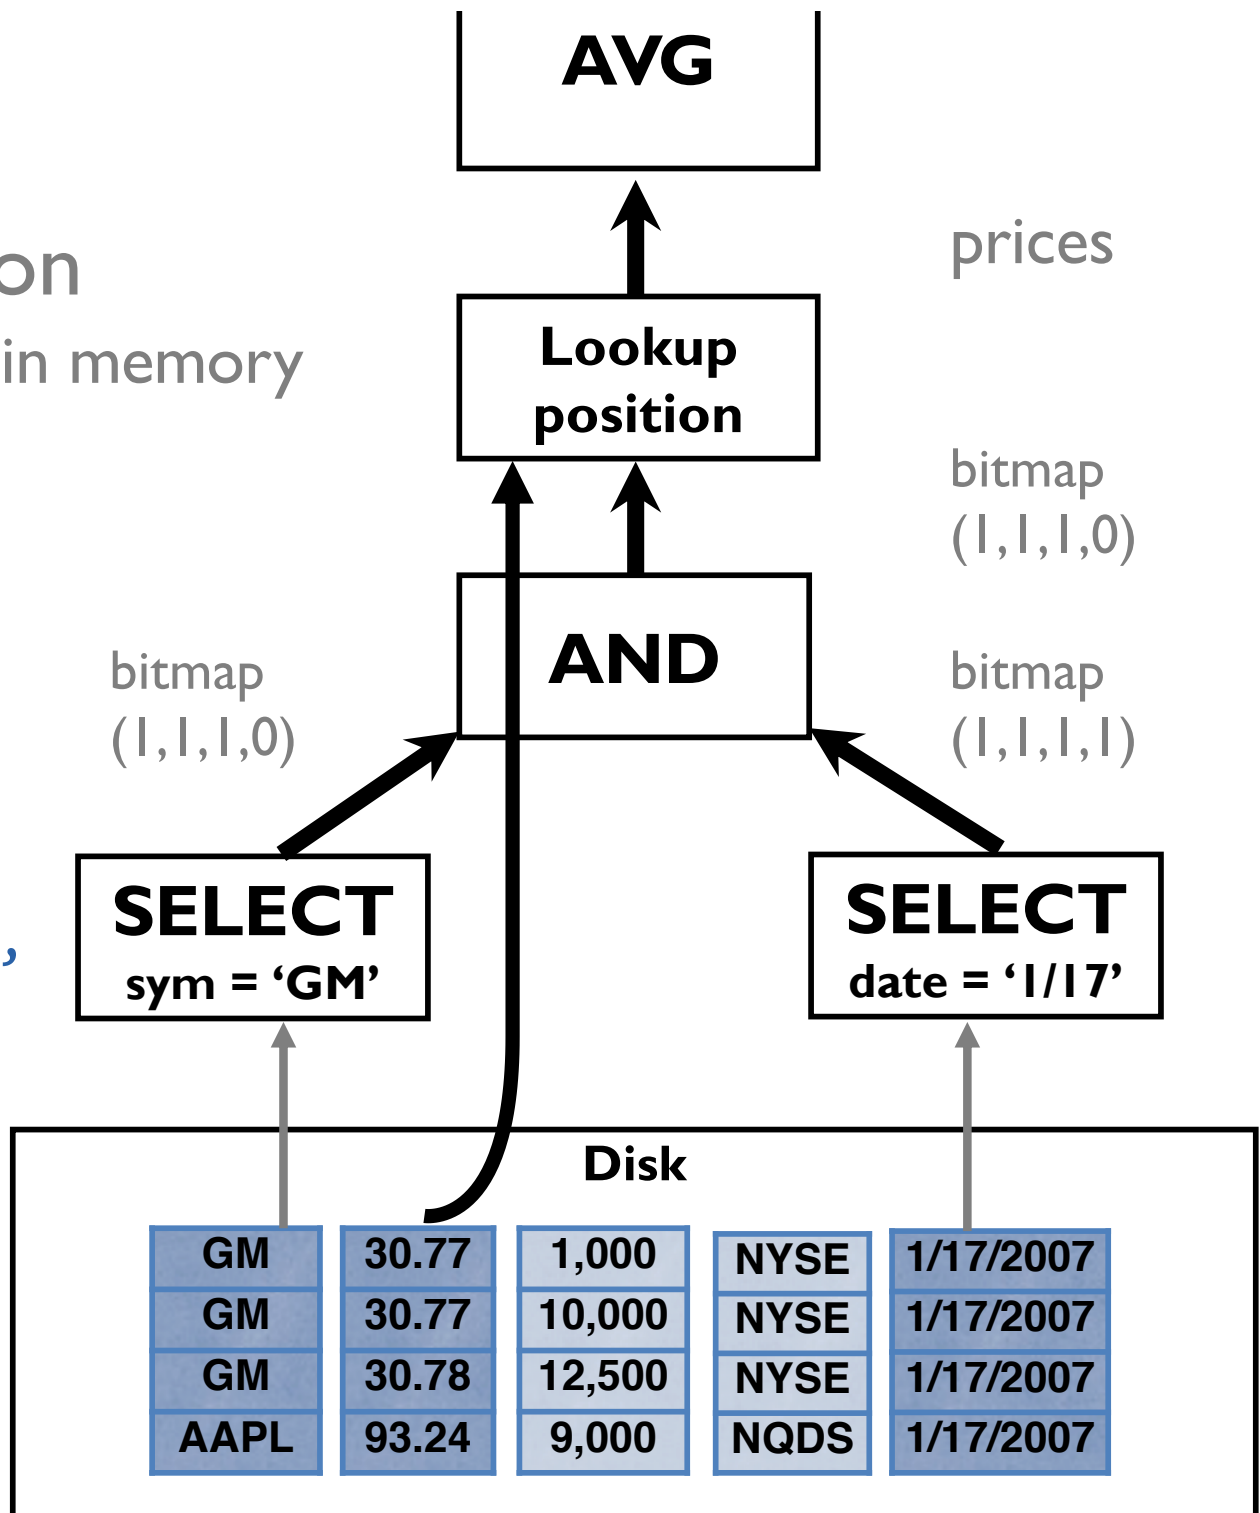


C-Store

Late Materialization

Much less data moving in memory

```
SELECT avg(price)
FROM tickstore
WHERE symbol = 'GM'
AND date = '1/17/2007'
```

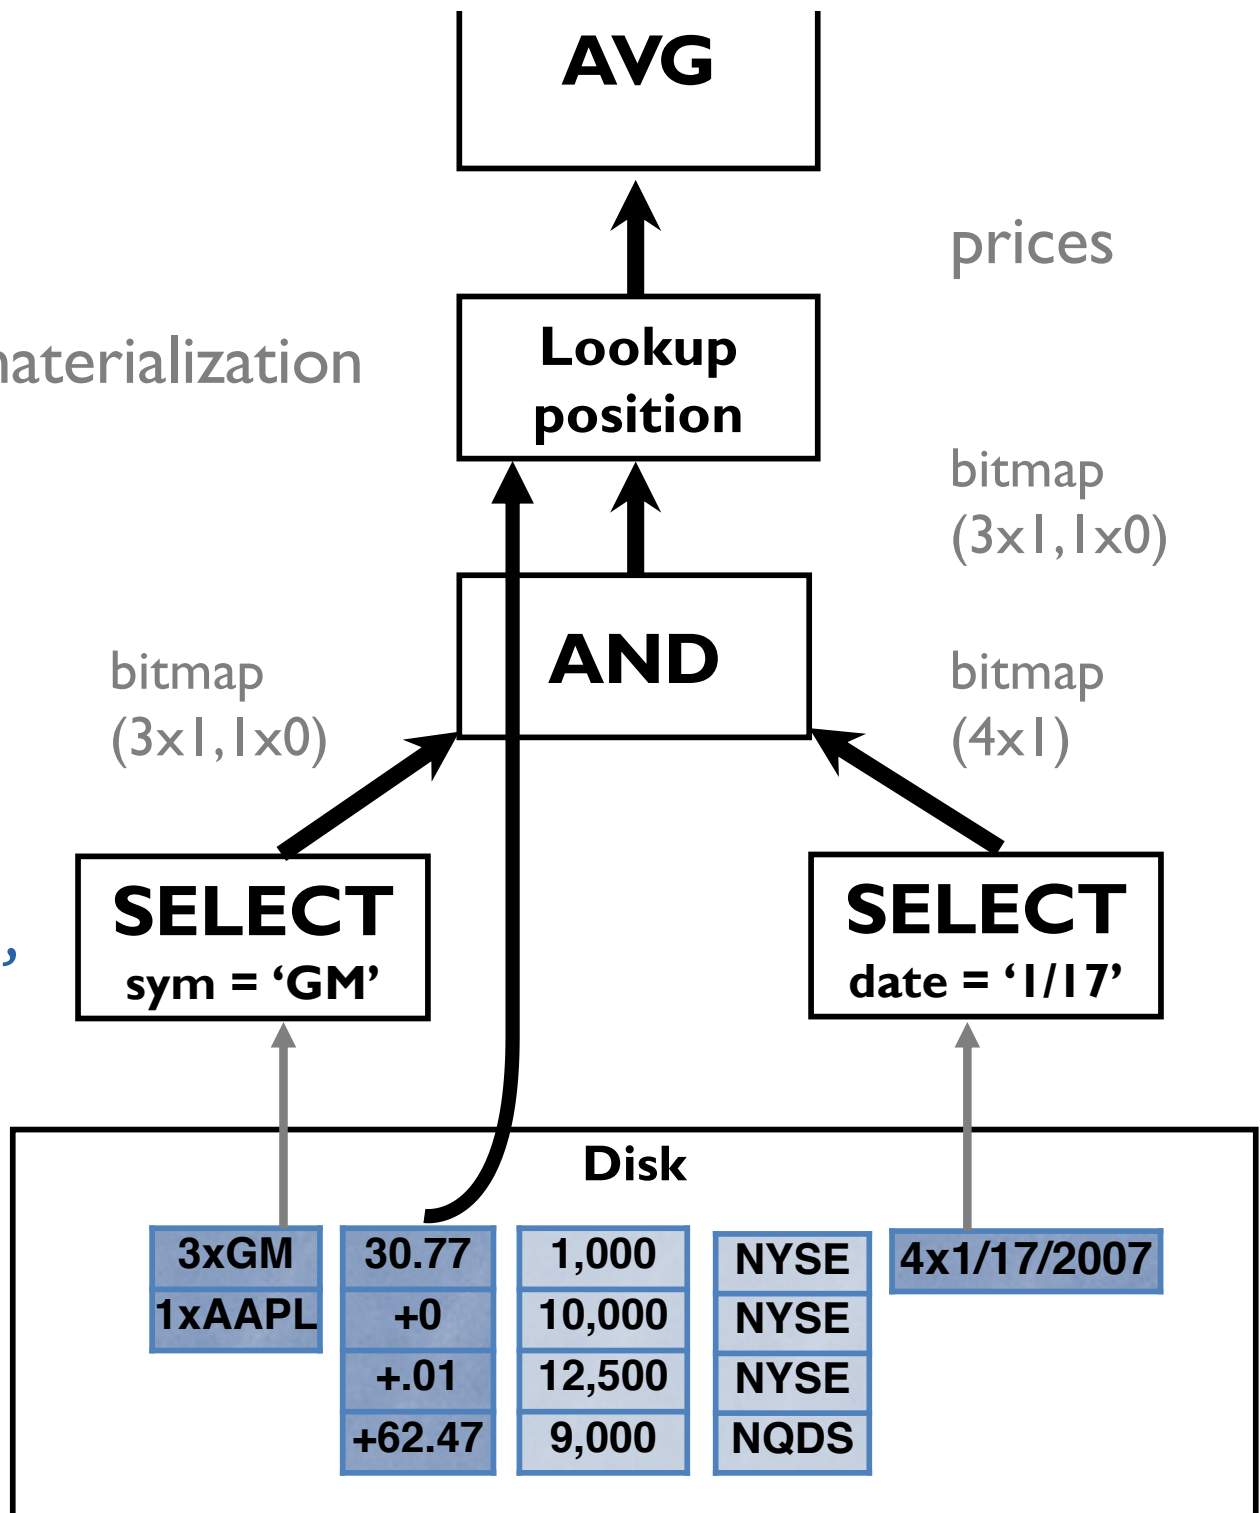


C-Store

Compression

Only possible w/ late materialization

```
SELECT avg(price)
FROM tickstore
WHERE symbol = 'GM'
AND date = '1/17/2007'
```



Column Stores

Optimized for data warehouses

Store data by attribute/column rather than row

Compression

Compressed *query plan execution*

50-100x faster than row store

In-Memory DBMSes

Transaction-oriented apps

remove 1 unit from product

move 5 units from org 1 to org 2

(shopping carts, inventory)

Data stored in memory

Disk only used for check pointing

No concurrency

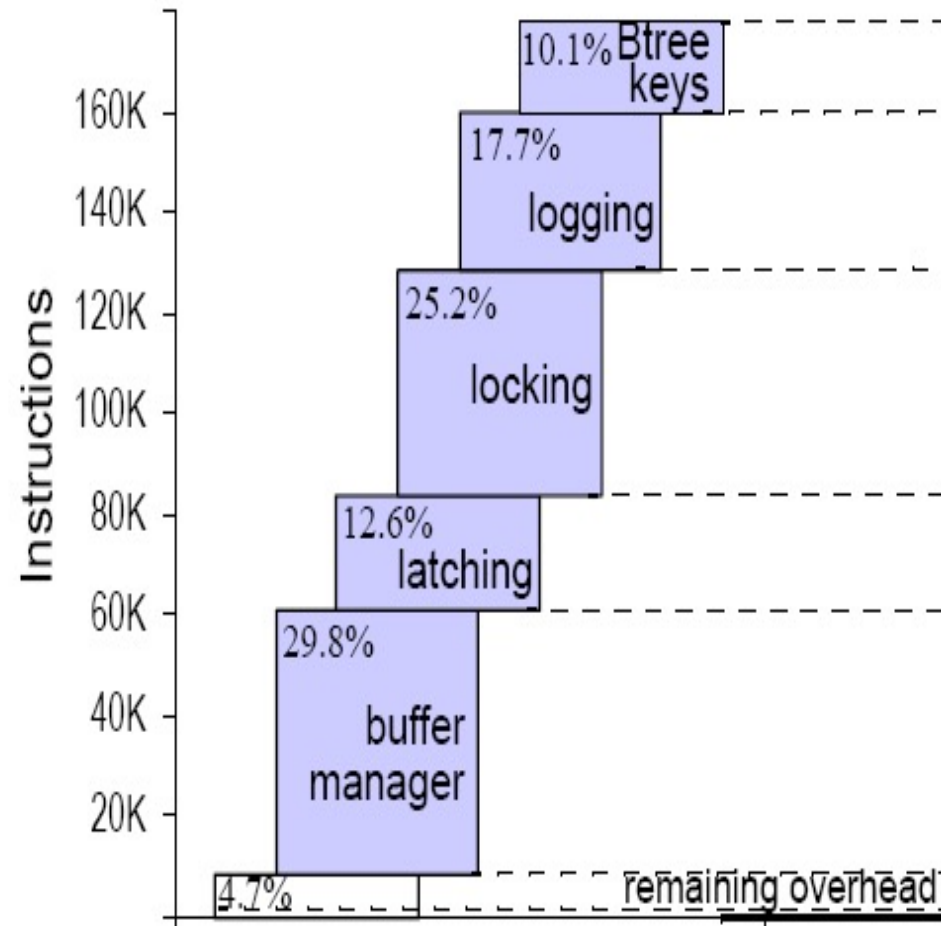
Active-active replication for fault-tolerance

Traditional Database


Indexes	queries go faster
Concurrency	queries go faster
Locking	serializability
Logging	recovery
Buffer Manager	manage pages in memory

Results after removing the components (in # instruction)


Instruction of useful work is only <2% of a memory resident DB




Effect of removing different components for payment (1/6)




Effect of removing different components for payment (2/6)




Effect of removing different components for payment (3/6)




Effect of removing different components for payment (4/6)



Effect of removing different components for payment (5/6)



Effect of removing different components for payment (6/6)





A diagram consisting of two vertically aligned components. The top component is a rectangle with a black border, containing the text "Barebones Executor". The bottom component is a cylinder with a black outline, containing the text "Disk".

Barebones
Executor

Disk

In memory DB



What about

Parsing

Concurrency?

Recovery?

In memory DB



What about

Parsing

Concurrency?

Recovery?

Procedure:

```
p1 = SELECT cost
      FROM fact_table
      WHERE id = ?
```

Query:

```
p1(10)
```

In memory DB



What about

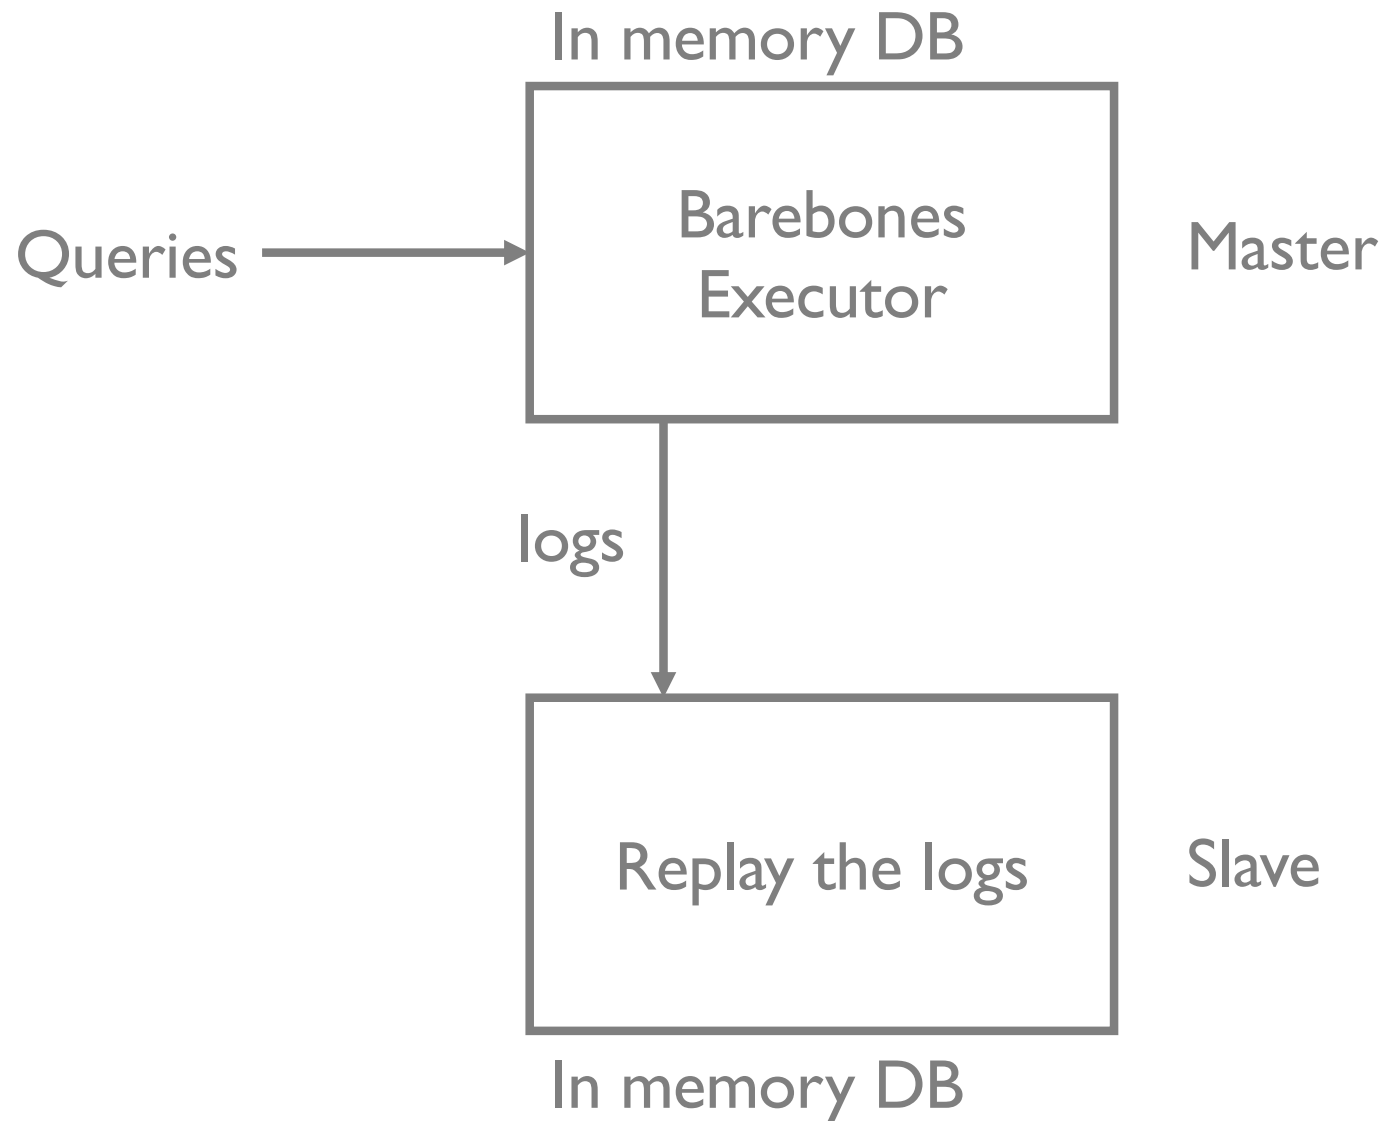
Parsing

~~Concurrency?~~


no buffer manager, no concurrency, no locks

Recovery?


Log Shipping




Active-Active




Recovery in Active-Active



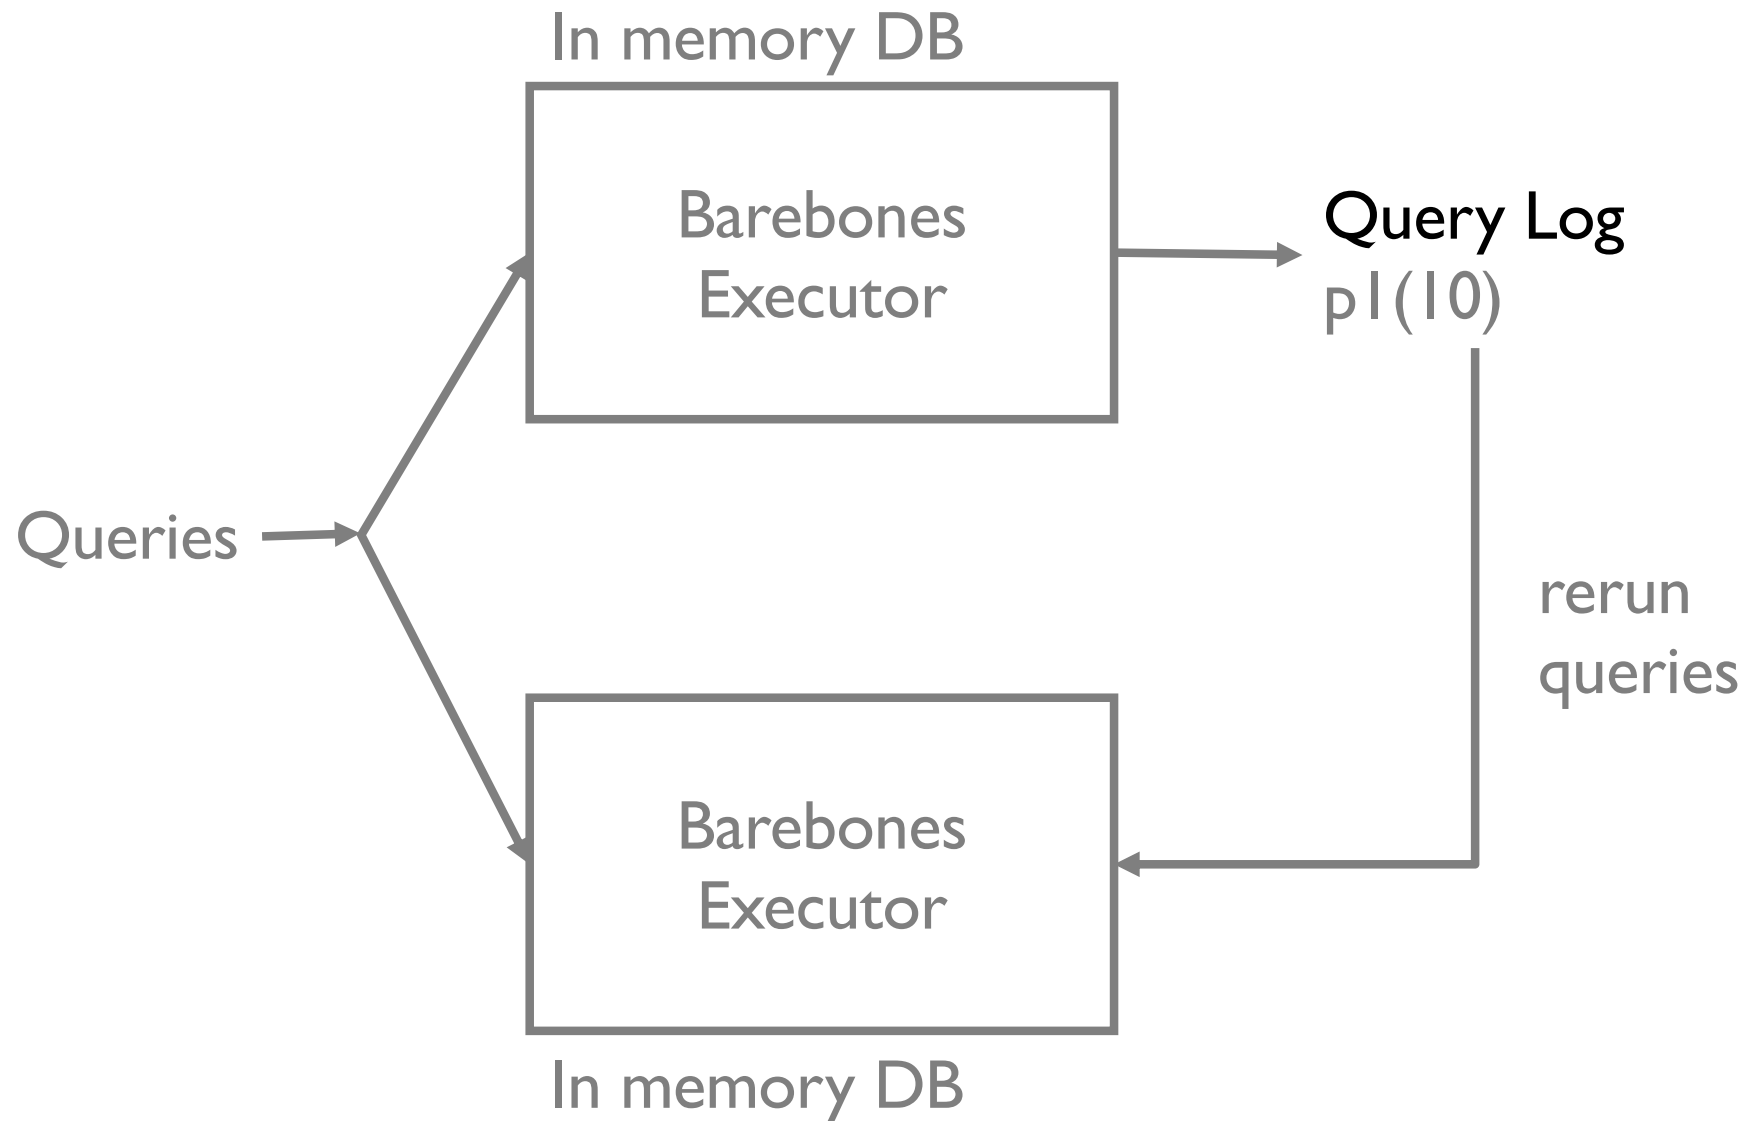
Recovery in Active-Active






Recovery in Active-Active




Recovery in Active-Active




One Size **Does Not** Fits All



One Size **Does Not** Fits All



One Size **Does Not** Fits All



OK Let's Step Back

Discussed the *how*:

- how to model data needed by an application

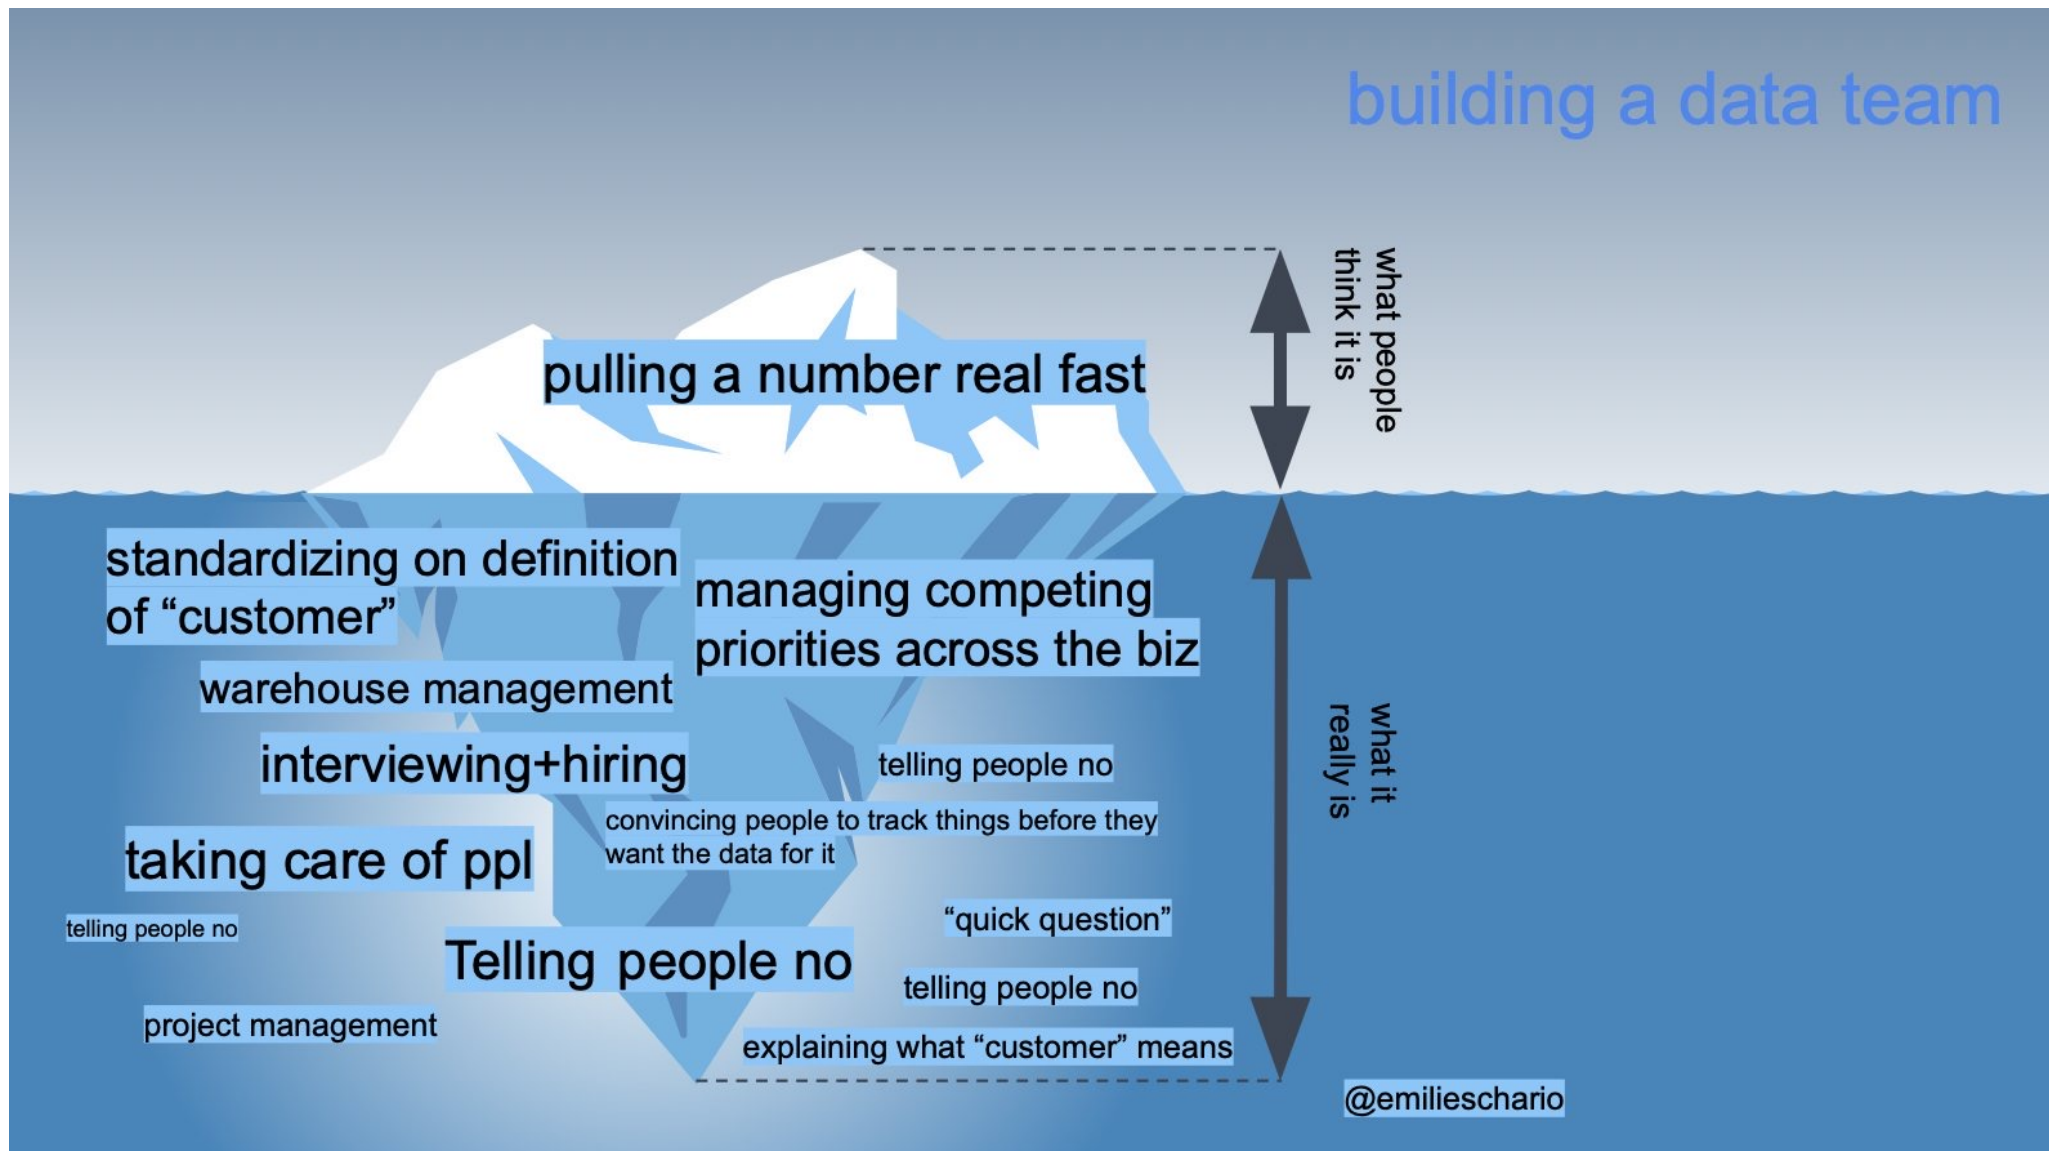
- how to query databases

- how databases execute queries

- how databases run fast, correctly

To what end?

building a data team



Production ML

An on-call engineer's biggest nightmare

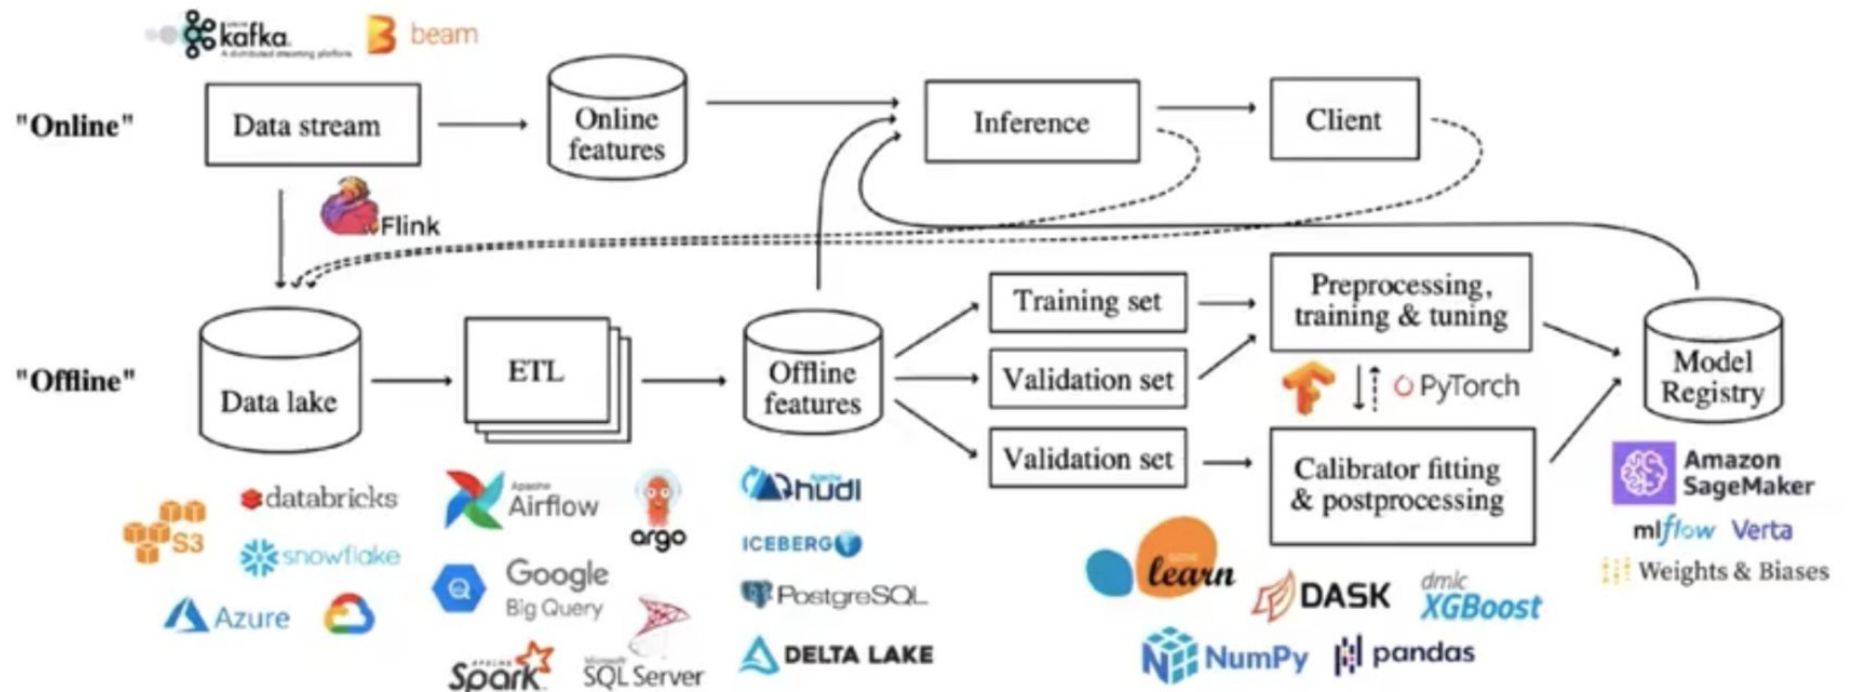


Figure 1: High-level architecture of a generic end-to-end machine learning pipeline. Logos represent a sample of tools used to construct components of the pipeline, illustrating heterogeneity in the tool stack. Shankar et al. 2021

<https://www.facebook.com/Engineering/videos/1578607659138164/>

4xxx: “here are facts about data management”

- 4111: basics
- 4112: gory DBMS internals

6xxx: “how people figured out those facts”

- 6111: info extraction and web
- 6113: class and modern DB research
- 6998: systems for human data interaction

PhD: “we don’t understand anything. Plz help”

Thanks!

Please fill out courseworks survey
cool.